

RL PAPER 3rd

APE-X

Distributed Prioritized Experience Replay

Published as a conference paper at **ICLR 2018**

2020.12.07 Mon

Jungyeon Lee

Published as a conference paper at ICLR 2018

DISTRIBUTED PRIORITIZED EXPERIENCE REPLAY

Dan Horgan

DeepMind

horgan@google.com

John Quan

DeepMind

johnquan@google.com

David Budden

DeepMind

budden@google.com

Gabriel Barth-Maron

DeepMind

gabrielbm@google.com

Matteo Hessel

DeepMind

mtthss@google.com

Hado van Hasselt

DeepMind

hado@google.com

David Silver

DeepMind

davidsilver@google.com

순서

INTRODUCTION

BACKGROUND

- 3 keywords

OUR CONTRIBUTION

- Structure
- pseudo code
- APE-X DQN
- APE-X DPG

EXPERIMENTS

- Atari
- Continuous control

ANALYSIS

- Actor 수에 대한 분석
- Replay memory에 대한 분석
- Scalability의 다른 요소들에 대한 분석

CONCLUSION

APPENDIX

Abstract

a **distributed** architecture for DRL at scale, that **enables agents to learn effectively from orders of magnitude more data** than previously possible.

- **decouple acting from learning**:
 - **the Actors**
 - interact with **their own** instances of the environment by selecting actions according to **a shared neural network**
 - accumulate the resulting experience **in a shared experience replay memory**
 - **the Learner**
 - replay samples of experience and updates the neural network
- rely on **prioritized experience replay** to focus only on the most significant data generated by the actors.
- substantially improve the state of the art on the Arcade Learning Environment, achieving better final performance in a fraction of the wall-clock training time.



Introduction

이 연구에 대해 빠르게 소개하겠습니다

딥러닝의 트렌드

- 더 파워풀한 모델 + 더 큰 데이터셋 = 더 좋은 결과
- effective use of greater computational resource 가 중요한 성공 요인

어떻게 scale up 할 것인가

- TF 프레임워크에 의한 scale up 이 아닌,
- **싱글 머신**에서 DRL 자체로 컴퓨팅을 개선 시킬 수 있을까?

이 연구에서는

Data Generation/Selection

- 기존: Gradient들을 병렬 계산해서 파라미터들을 빨리 업데이트

1. distribute the generation

2. selection of experience data

- 기존 방법들의 대체제
- 2가지 방법을 결합할 수 있지만 (1) 데이터 생산에 초점

실험

- DQN(Deep Q-Networks)
- DDPG(Deep Deterministic Policy Gradient)
- Arcade Learning Env
- Continuous control tasks
- 하이퍼파라미터 튜닝 안하고 SOTA

분석

- how prioritization affects performance as we increase the number of data-generating workers
- replay capacity
- recency of the experience
- use of different data-generating policies for different workers

Background

3가지 키워드로 보는 배경 지식



Distributed SGD

- @ 지도학습 학습시간을 가속하기 위해 흔히 쓰는 방법
- 그래디언트 계산을 병렬화
- 파라미터 업데이트는 동시성/비동시성 둘 다 가능
- 비동시성 파라미터 업데이트+분산 데이터 생산 : 멀티 쓰레드로 성공적이었음
- GA3C/PAAC 등등



Distributed Importance Sampling

- 학습 빠르게 하기
 - IS로 variance ▼ bias ▲
 - 수렴 속도는 증가
 - gradient의 L2 norm에 비례하여 샘플링 → @ 지도학습 성공적
1. Rank samples according to their *latest known loss value*
 2. Make the sampling probability *a function of the rank* rather than of the loss itself




Prioritized Experience Replay

- **Prioritized** : 데이터 효율성 증대 (training NN이 SGD 알고리즘을 사용할 때 효율적)
- **Experience Replay** : 오버피팅을 방지
- 좀 더 *general*한 biased sampling (가장 surprising한 경험을 배우는 데에 초점)
- Rainbow에서 봤듯이 그 효과는 대단!

Distributed PER

EXTEND PER TO THE DISTRIBUTED SETTING

- @ Gorila
 - **acting :**
 - stepping through an environment
 - evaluating a policy implemented as a deep neural network
 - storing the observed data in a replay memory
 - **learning:**
 - sampling batches of data from the memory to update the policy parameters
- @ APE-X
 - actor N개 : CPU
 - learner 1개 : GPU

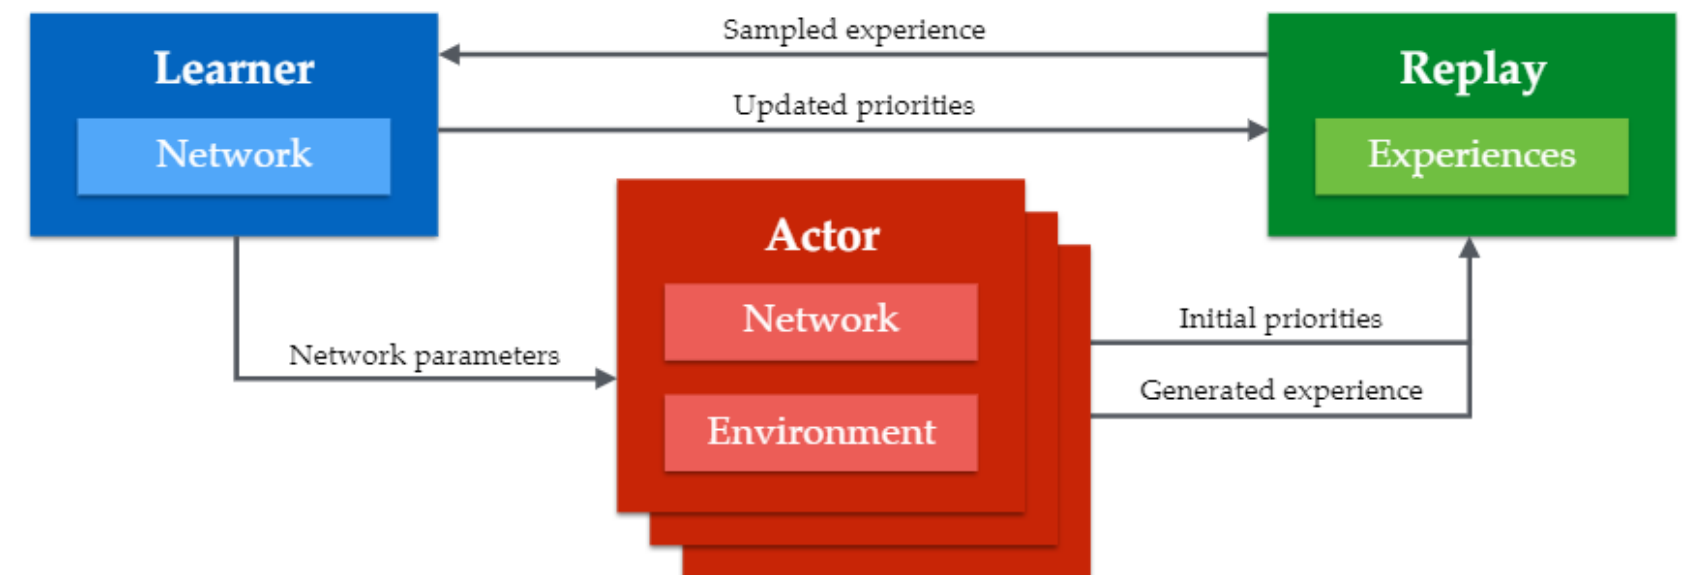


Distributed PER

EXTEND PER TO THE DISTRIBUTED SETTING

- Shared, Centralized **replay memory**
- **Prioritize** to sample the most useful data more often
 - shared(어떤 actor든지 발견하기만 하면 됨)
- **Priority Definition**
 - a. @ Prioritized DQN : sampling된 sample들만 priority를 update / scale up 잘 안됨

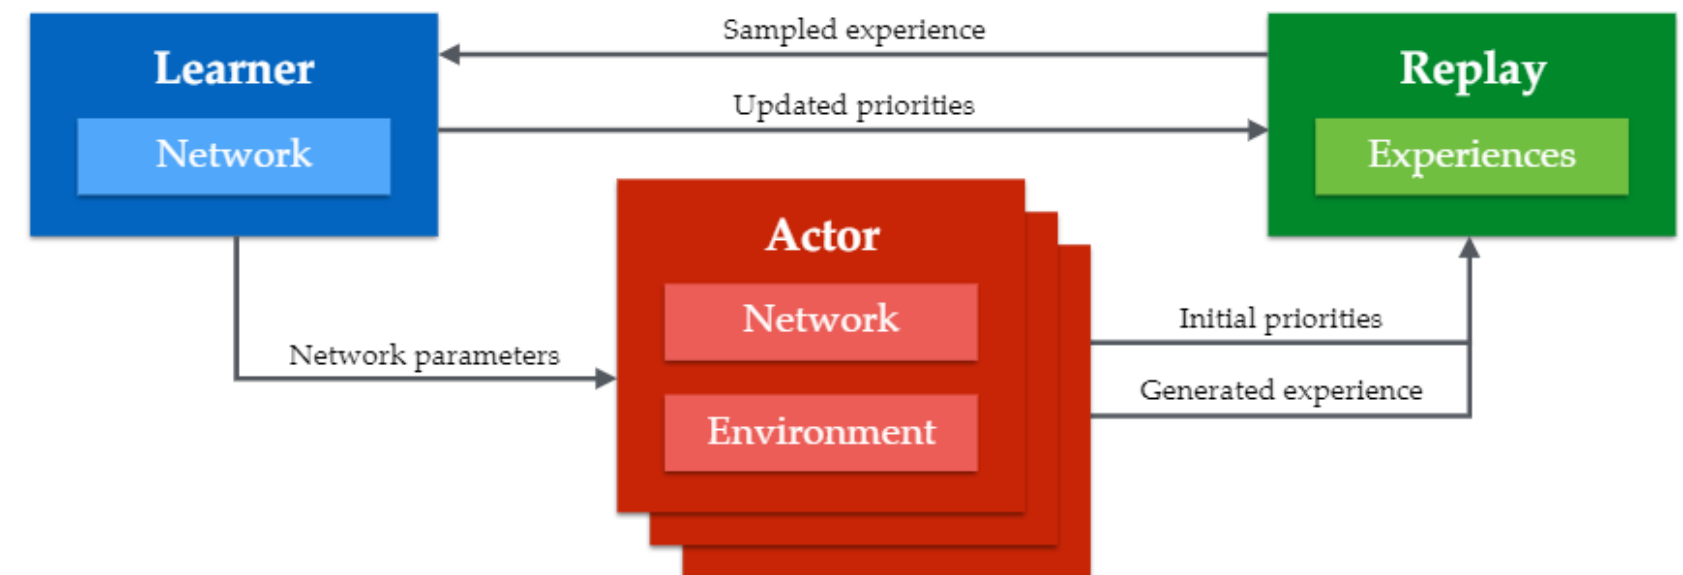
@ **APE-X** : actor들이 replay에 넣을 때 priority를 다시 계산하니, 추가적인 computation없이 큰 문제를 해결

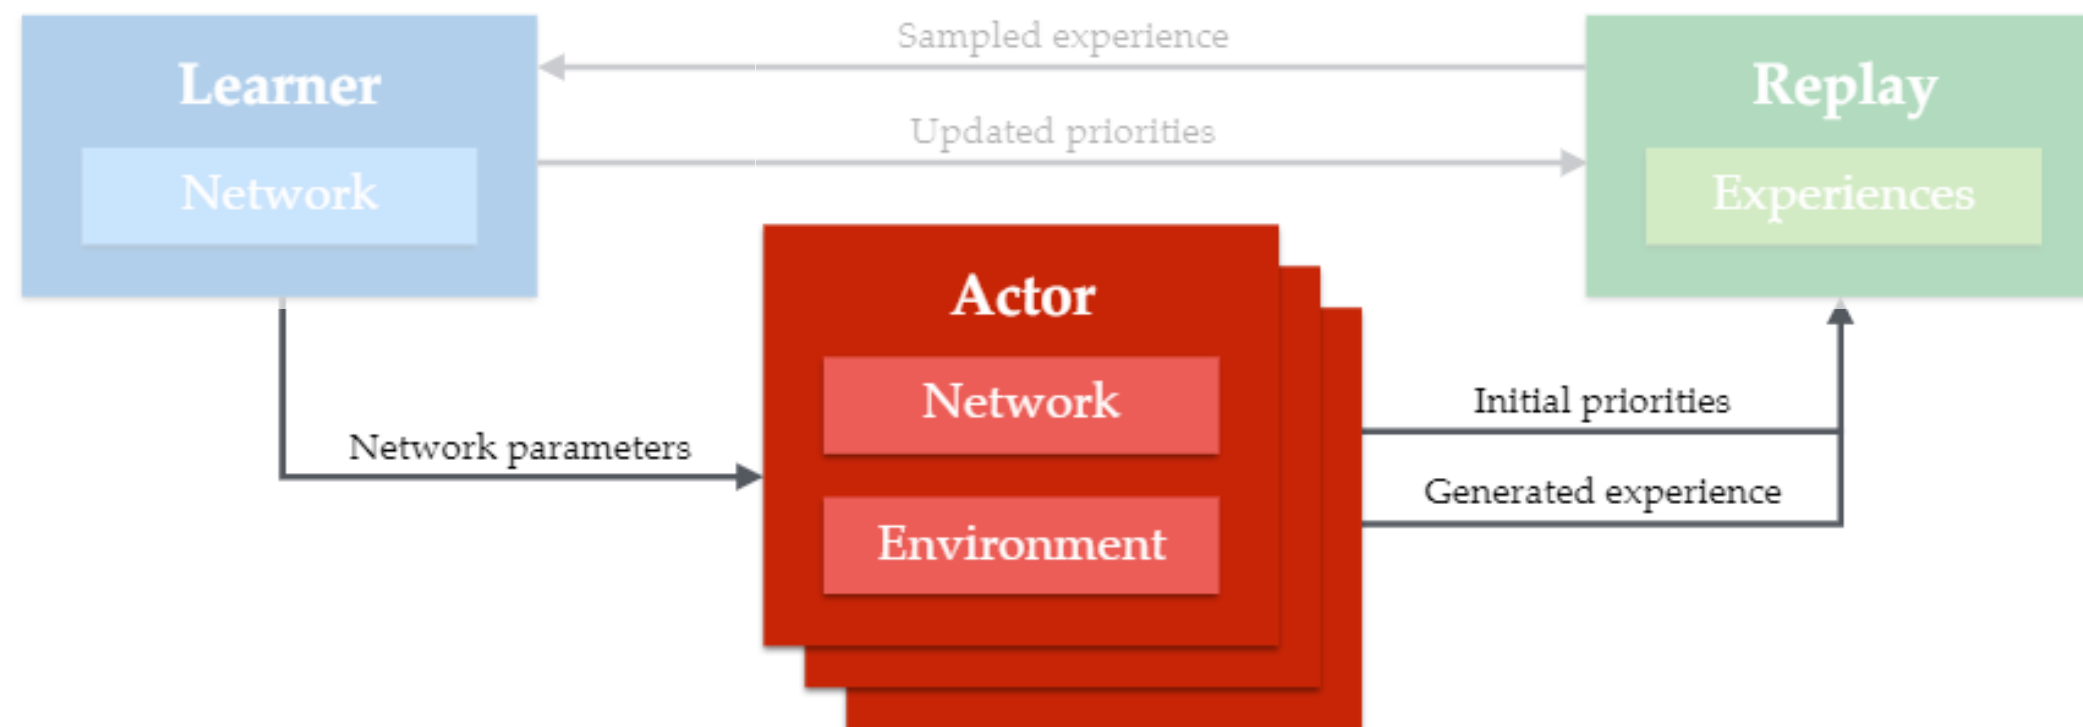


Distributed PER

EXTEND PER TO THE DISTRIBUTED SETTING

- **Sharing experiences** VS Sharing gradients
 - 경험 데이터의 outdated가 그라디언트보다 늦음
→ off policy 데이터에 대해 robust
 - centralized replay
 - actor와 learner들이 다른 데이터 센터에서 돌아갈 수 있음(퍼포먼스에 제한이 없음)
- **Off policy**
 - 각각의 actor들이 다른 policy로 다양하게 경험
→ exploration에 대해서 진전이 있었음





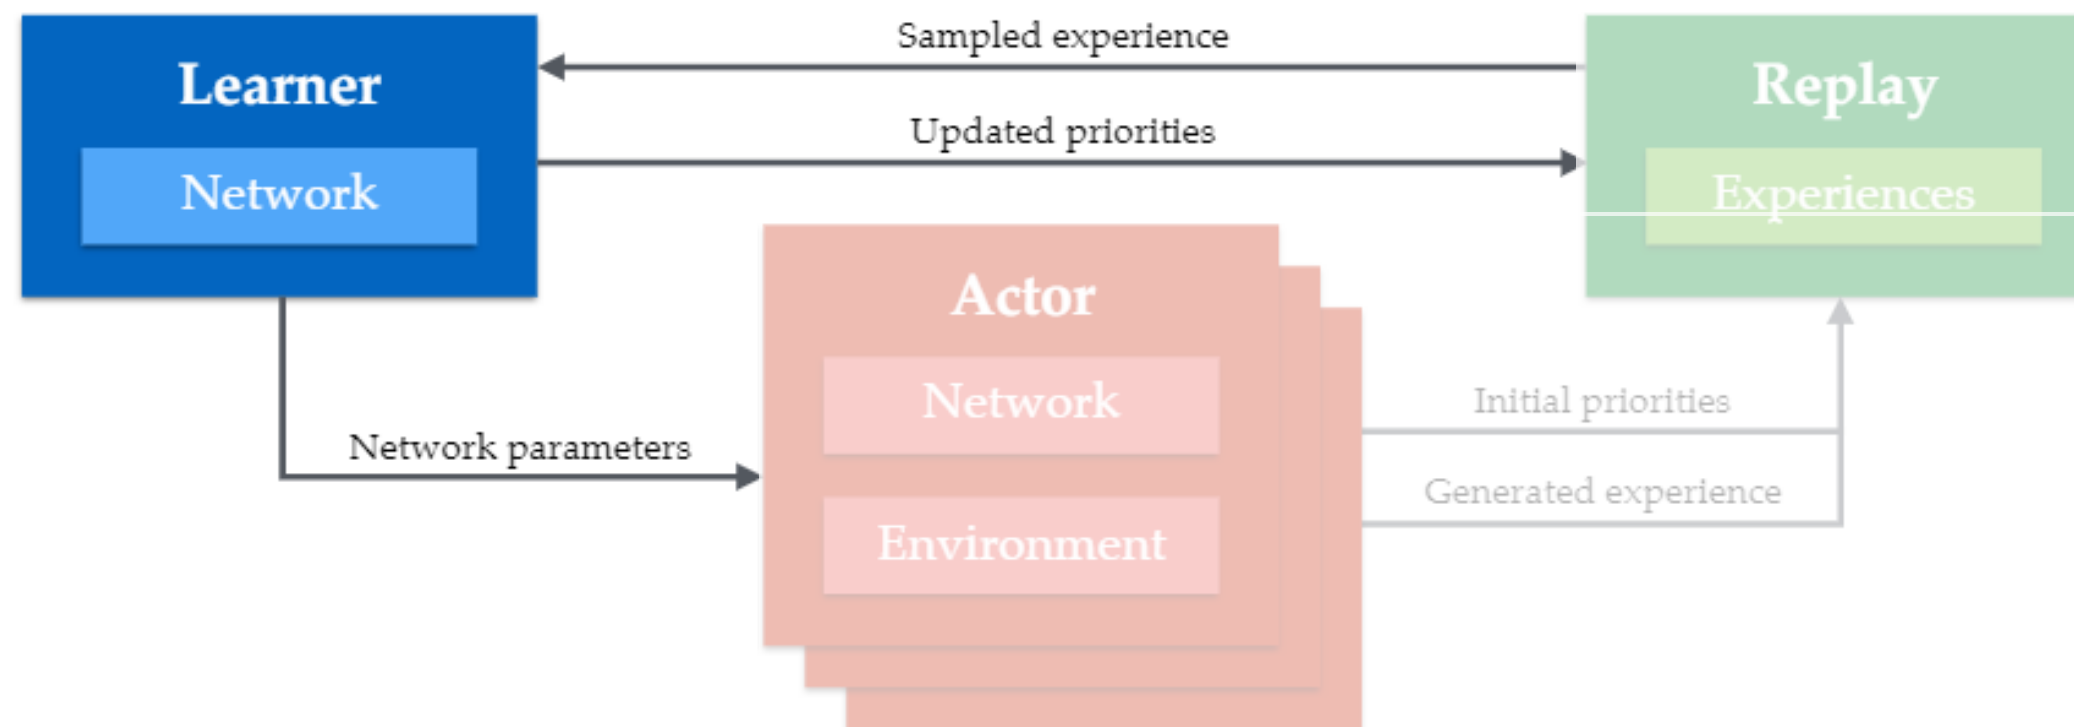
Algorithm 1 Actor

```

1: procedure ACTOR( $B, T$ )
2:    $\theta_0 \leftarrow \text{LEARNER.PARAMETERS}()$ 
3:    $s_0 \leftarrow \text{ENVIRONMENT.INITIALIZE}()$ 
4:   for  $t = 1$  to  $T$  do
5:      $a_{t-1} \leftarrow \pi_{\theta_{t-1}}(s_{t-1})$ 
6:      $(r_t, \gamma_t, s_t) \leftarrow \text{ENVIRONMENT.STEP}(a_{t-1})$ 
7:     LOCALBUFFER.ADD $((s_{t-1}, a_{t-1}, r_t, \gamma_t))$ 
8:     if LOCALBUFFER.SIZE $() \geq B$  then
9:        $\tau \leftarrow \text{LOCALBUFFER.GET}(B)$ 
10:       $p \leftarrow \text{COMPUTE PRIORITIES}(\tau)$ 
11:      REPLAY.ADD $(\tau, p)$ 
12:     end if
13:     PERIODICALLY $(\theta_t \leftarrow \text{LEARNER.PARAMETERS}())$ 
14:   end for
15: end procedure

```

▷ Run agent in environment instance, storing experiences.
 ▷ Remote call to obtain latest network parameters.
 ▷ Get initial state from environment.
 ▷ Select an action using the current policy.
 ▷ Apply the action in the environment.
 ▷ Add data to local buffer.
 ▷ In a background thread, periodically send data to replay.
 ▷ Get buffered data (e.g. batch of multi-step transitions).
 ▷ Calculate priorities for experience (e.g. absolute TD error).
 ▷ Remote call to add experience to replay memory.
 ▷ Obtain latest network parameters.



Algorithm 2 Learner

```

1: procedure LEARNER( $T$ )                                ▷ Update network using batches sampled from memory.
2:    $\theta_0 \leftarrow \text{INITIALIZE\_NETWORK}()$ 
3:   for  $t = 1$  to  $T$  do                                  ▷ Update the parameters  $T$  times.
4:      $id, \tau \leftarrow \text{REPLAY.SAMPLE}()$                 ▷ Sample a prioritized batch of transitions (in a background thread).
5:      $l_t \leftarrow \text{COMPUTE\_LOSS}(\tau; \theta_t)$           ▷ Apply learning rule; e.g. double Q-learning or DDPG
6:      $\theta_{t+1} \leftarrow \text{UPDATE\_PARAMETERS}(l_t; \theta_t)$ 
7:      $p \leftarrow \text{COMPUTE\_PRIORITIES}()$                 ▷ Calculate priorities for experience, (e.g. absolute TD error).
8:      $\text{REPLAY.SET\_PRIORITY}(id, p)$                       ▷ Remote call to update priorities.
9:      $\text{PERIODICALLY}(\text{REPLAY.REMOVE\_TO\_FIT}())$           ▷ Remove old experience from replay memory.
10:  end for
11: end procedure
  
```

APE-X DQN

ATARI



DQN

- learning algorithm : **Double DQN** with **multi-step** bootstrap targets
- **function approximator** : dueling network $q(\cdot, \cdot, \theta)$

→ batch Loss

$$l_t(\theta) = \frac{1}{2} (G_t - q(S_t, A_t, \theta))^2$$

→ Gain

$$G_t = \underbrace{R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n}}_{\text{multi-step return}} + \underbrace{\gamma^n \overbrace{q(S_{t+n}, \arg\max_a q(S_{t+n}, a, \theta), \theta^-)}^{\text{double-Q bootstrap value}}}_{\text{target network}}$$

APE-X DQN

ATARI

more details

- **Multi-step return with no off-policy correction**
 - adversely affect the value estimation
- Actor executes a **different** policy (off policy)
 - Experiences to be generated from a variety of strategies
 - Behaviour policy affects
 - exploration
 - the quality of function approximation
- Actors use **ϵ -greedy policies** with different values of ϵ
 - **Low** ϵ policies allow *exploring deeper* in the environment
 - **High** ϵ policies *prevent over-specialization*
 - constant during training

$$\epsilon_i = \epsilon^{1 + \frac{i}{N-1} \alpha}$$

$$\epsilon = 0.4, \alpha = 7$$

Ape-X DPG

CONTINUOUS CONTROL TASK

Continuous-action policy gradient system based on DDPG

- similar to Ape-X DQN
- the policy : φ
- Q-network : ψ
- The two networks are optimized separately, by minimizing different losses on the sampled experience

$$G_t = \underbrace{R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n}}_{\text{multi-step return}} + \underbrace{\gamma^n \overbrace{q(S_{t+n}, \underset{a}{\operatorname{argmax}} q(S_{t+n}, a, \boldsymbol{\theta}), \boldsymbol{\theta}^-)}^{\text{double-Q bootstrap value}}}_{\text{double-Q bootstrap value}}$$

$$G_t = \underbrace{R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n q(S_{t+n}, \pi(S_{t+n}, \phi^-), \psi^-)}_{\text{multi-step return}}$$

Ape-X DPG

CONTINUOUS CONTROL TASK

$$G_t = \underbrace{R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n q(S_{t+n}, \pi(S_{t+n}, \phi^-), \psi^-)}_{\text{multi-step return}}$$

more details

• Q-network

- output : an action-value estimate $q(s, a, \psi)$
- updated using TD learning with a multi-step bootstrap target $l_t(\psi) = \frac{1}{2} (G_t - q(S_t, A_t, \psi))^2$

→ Loss

• Policy network

- output : an action $A_t = \pi(S_t, \phi) \in \mathbb{R}^m$
- updated using policy **gradient ascent** on the estimated Q-value

→ Gradient

$$\nabla_{\phi} q(S_t, \pi(S_t, \phi), \psi)$$

$$A_t = \pi(S_t, \phi) \rightarrow \text{input to critic network}$$

Experiments

ATARI

APE-X DQN

- **360** actor machines / 1 CPU core
- 139 FPS each → Total ~50k(= 139X360) FPS = ~12.5K transitions (4개 frame 당 1개 action)
- **Actor** :
 - batch experience data **locally** before sending it to the replay
 - copy Learner's parameters every 400 frames(~2.8 sec)
- **Learner** :
 - 비동기적으로 512 transition을 16개의 batch로 가져옴
 - 1초 당 19 batch update = ~9.7K transitions에 대한 gradient들을 계산
- Observation : PNG codec / Episode length : limited to 50000 frames during training
- Capacity of the shared experience replay memory : 약 2 million transitions (2035050)
- Data is sampled according to **proportional prioritization**
 - a priority exponent of 0.6
 - an importance sampling exponent 0.4



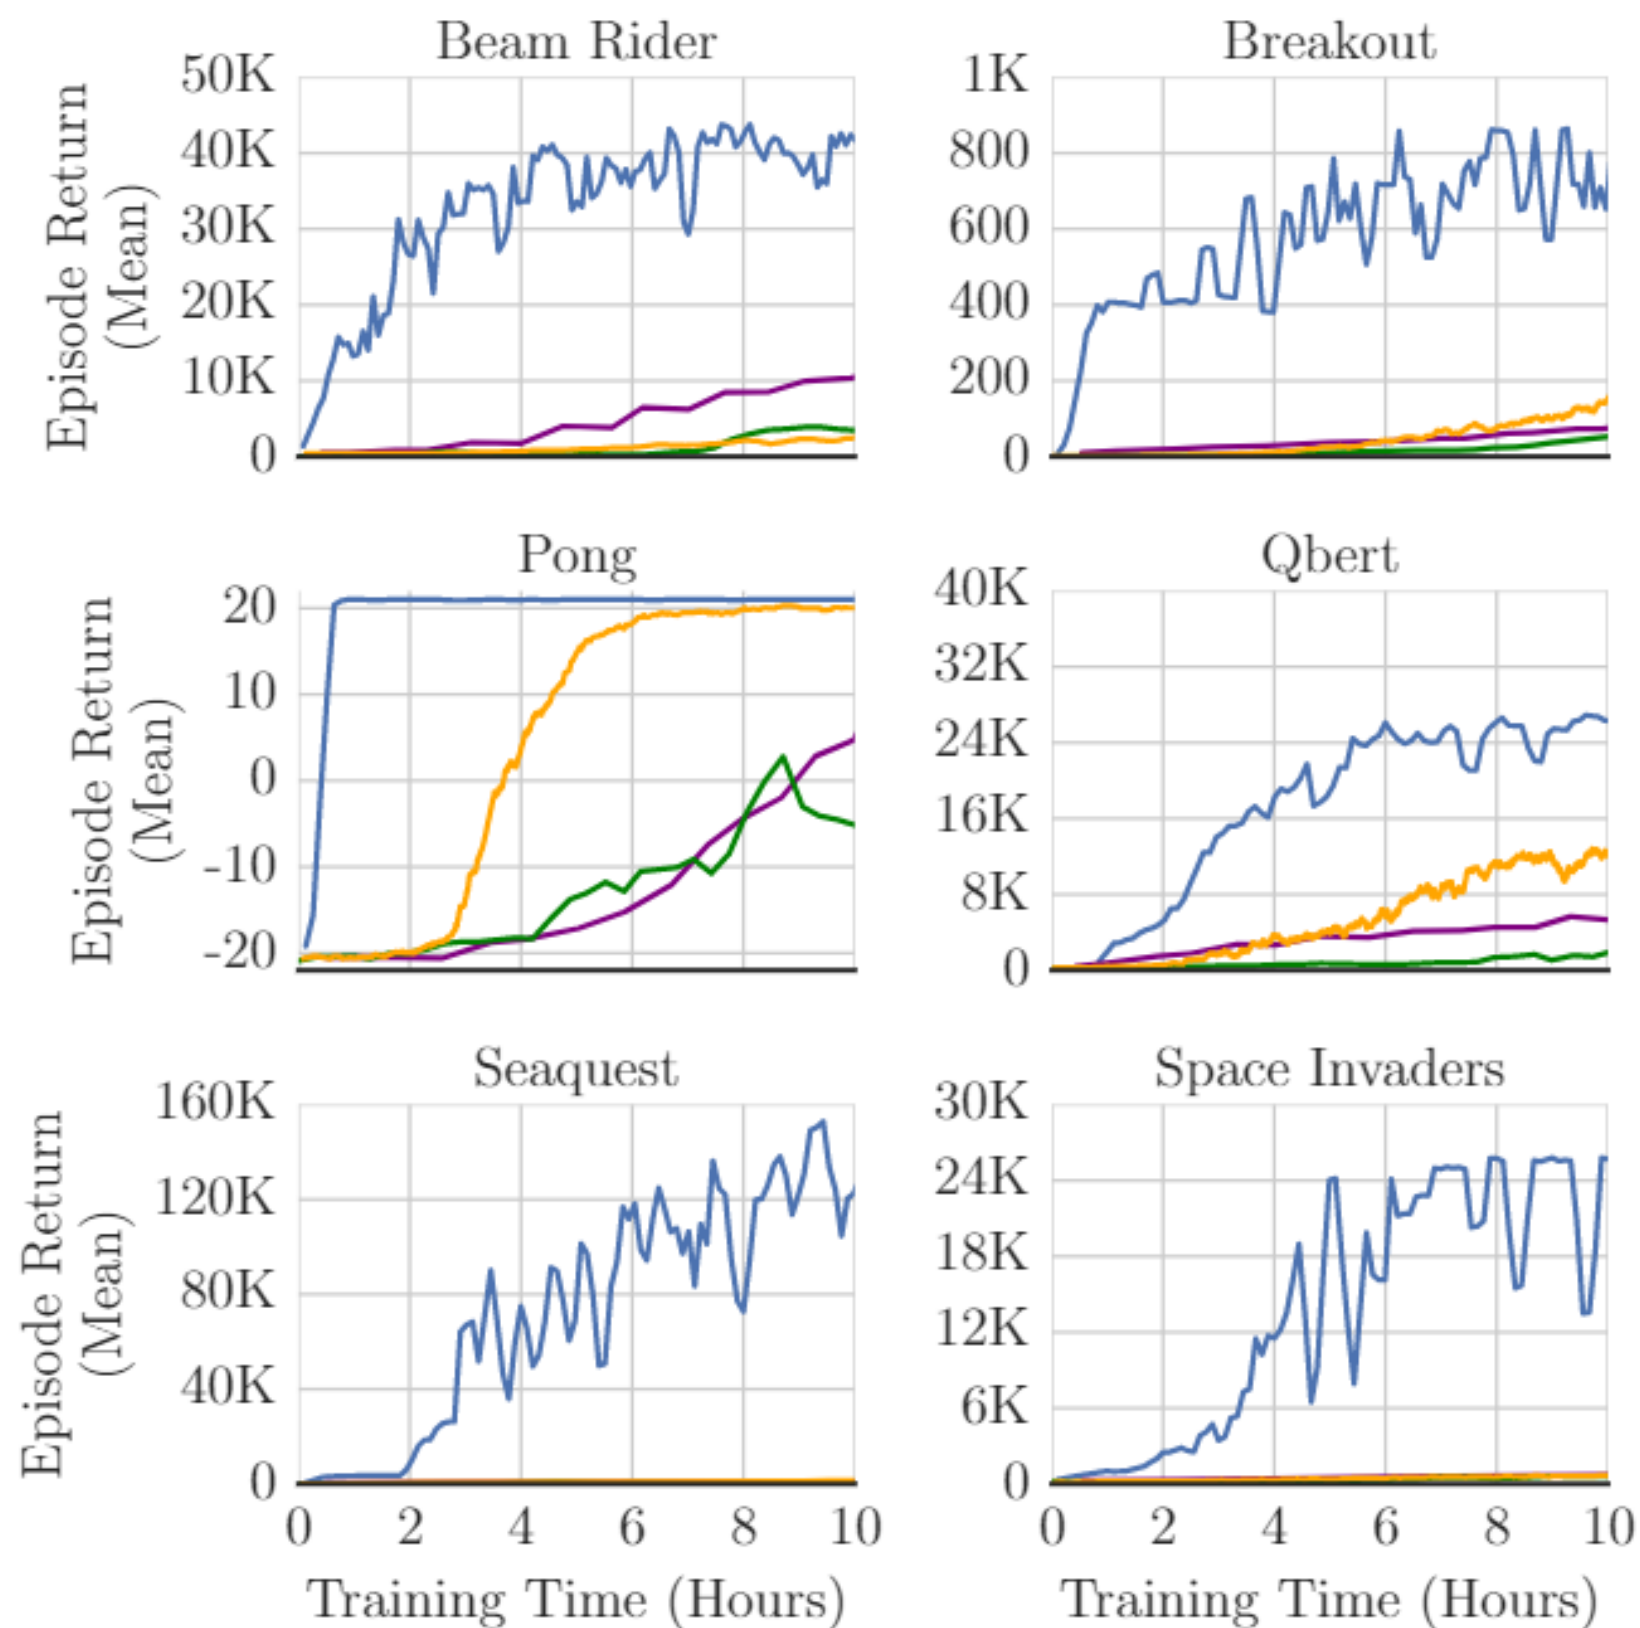



Figure 2: Left: Atari results aggregated across 57 games, evaluated from random no-op starts. Right: Atari training curves for selected games, against baselines. Blue: Ape-X DQN with 360 actors; Orange: A3C; Purple: Rainbow; Green: DQN. See appendix for longer runs over all games.

Experiments

ATARI

| Algorithm | Training Time | Environment Frames | Resources (per game) | Median (no-op starts) | Median (human starts) |
|-------------------------|---------------|--------------------|-------------------------------|-----------------------|-----------------------|
| Ape-X DQN | 5 days | 22800M | 376 cores, 1 GPU ^a | 434% | 358% |
| Rainbow | 10 days | 200M | 1 GPU | 223% | 153% |
| Distributional (C51) | 10 days | 200M | 1 GPU | 178% | 125% |
| A3C | 4 days | — | 16 cores | — | 117% |
| Prioritized Dueling | 9.5 days | 200M | 1 GPU | 172% | 115% |
| DQN | 9.5 days | 200M | 1 GPU | 79% | 68% |
| Gorila DQN ^c | ~4 days | — | unknown ^b | 96% | 78% |
| UNREAL ^d | — | 250M | 16 cores | 331% ^d | 250% ^d |

Table 1: Median normalized scores across 57 Atari games. ^a Tesla P100. ^b >100 CPUs, with a mixed number of cores per CPU machine. ^c Only evaluated on 49 games. ^d Hyper-parameters were tuned per game.

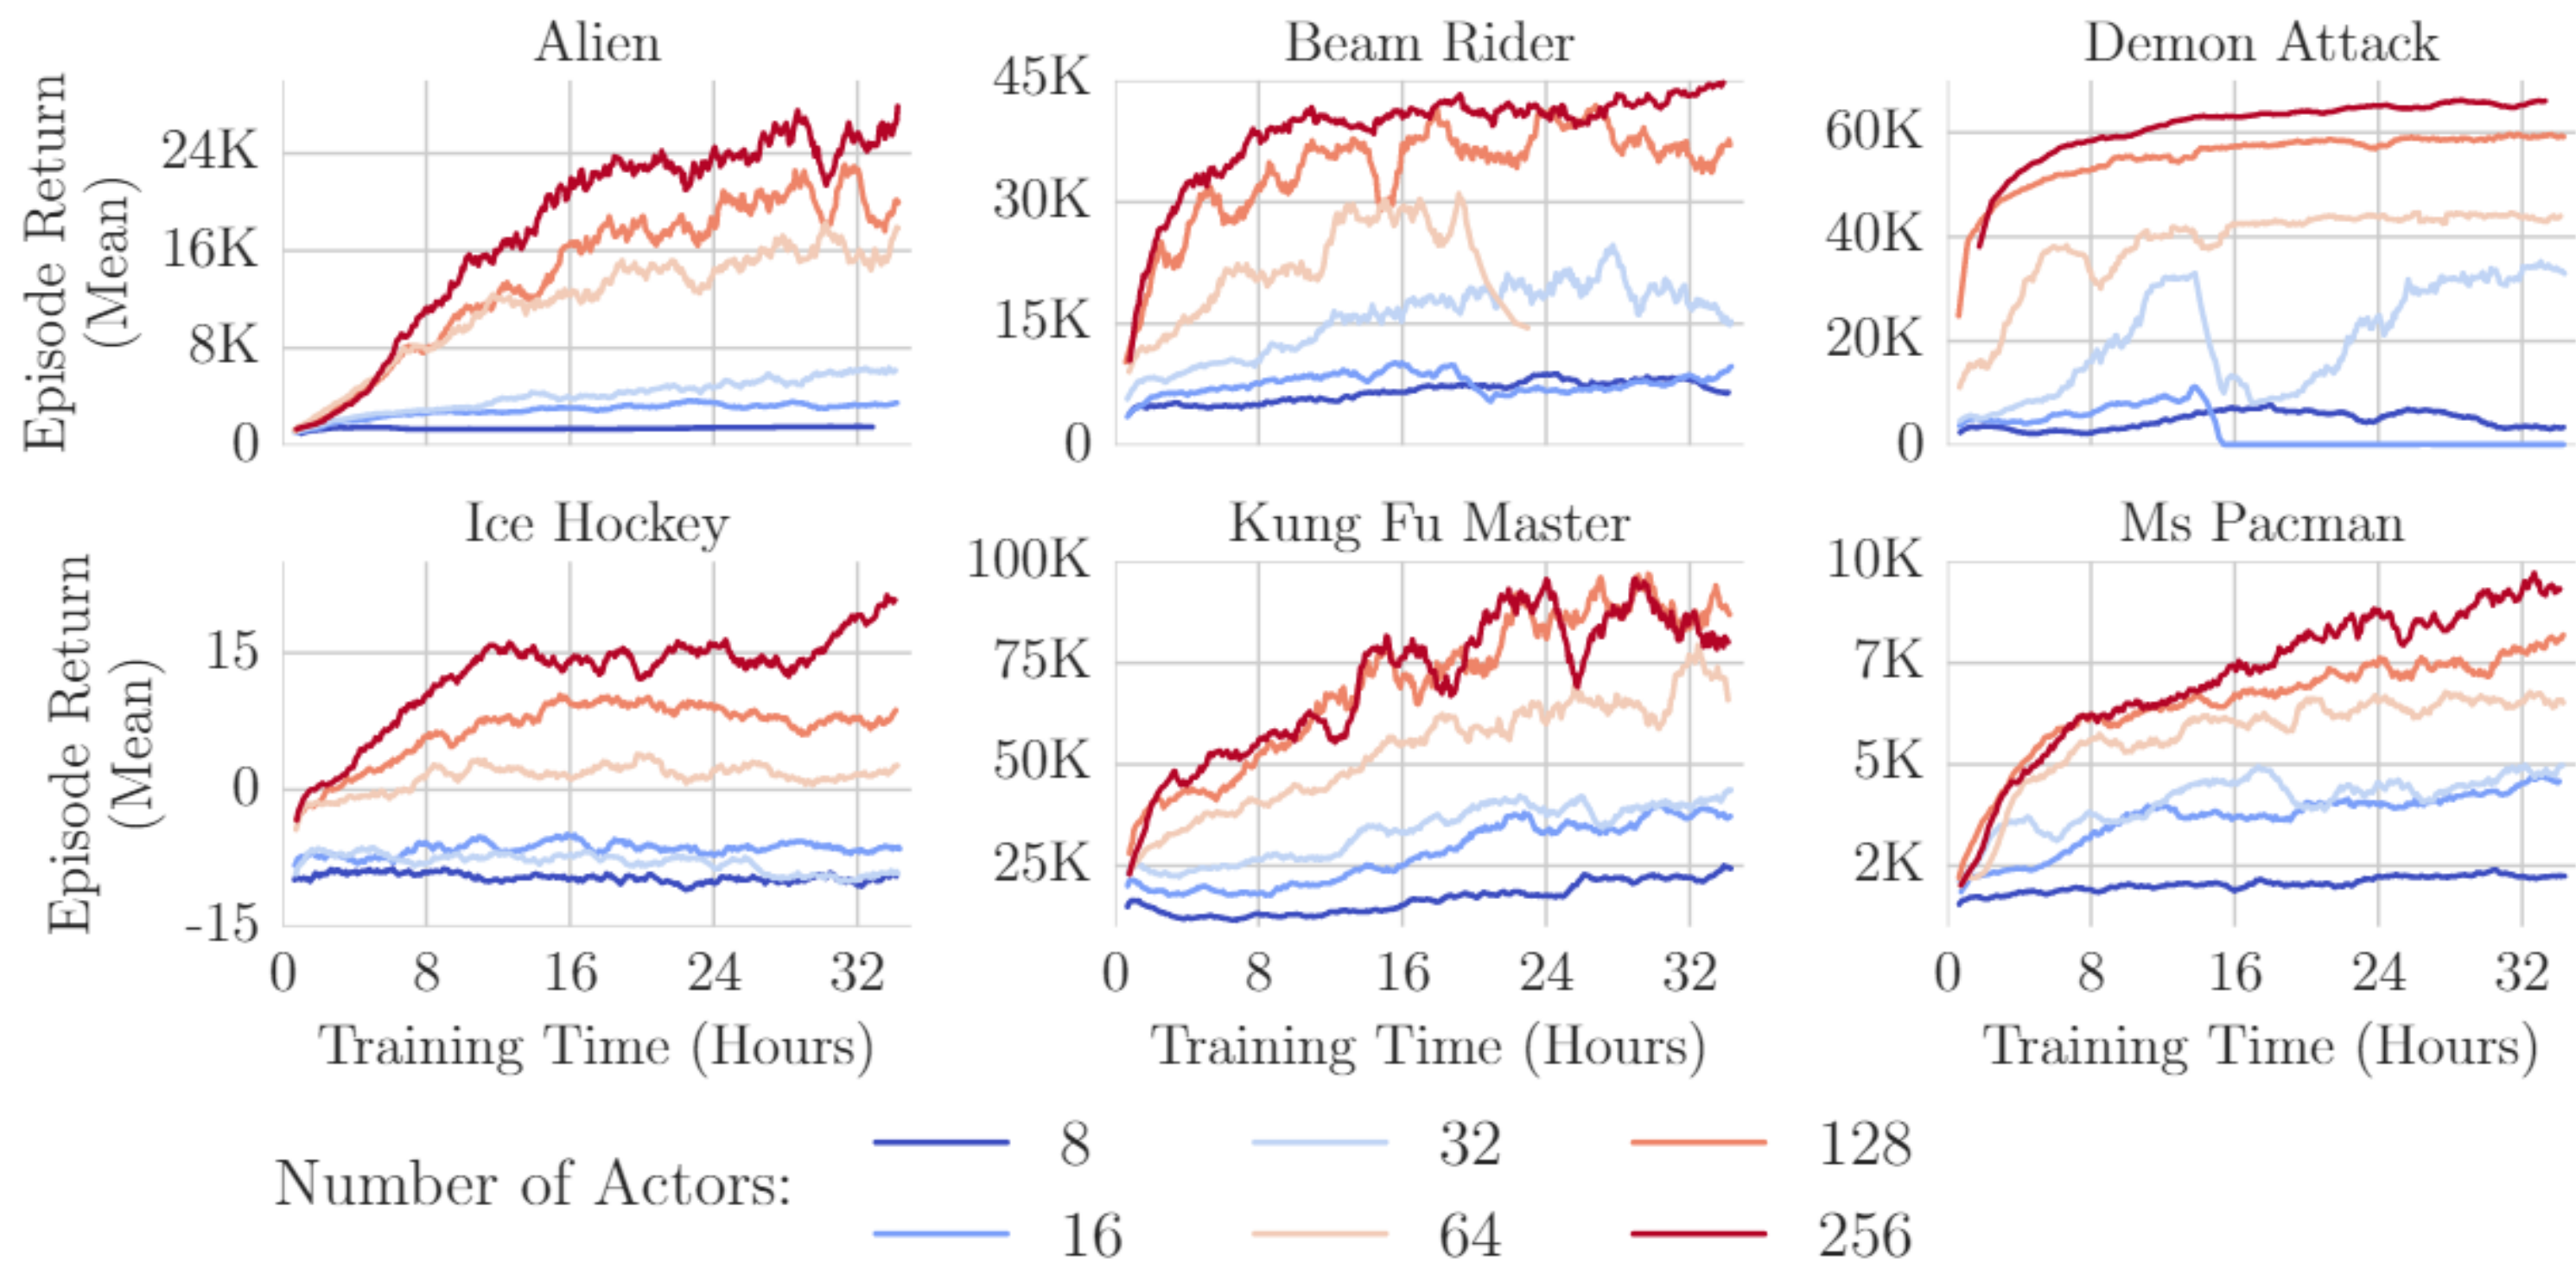


Figure 4: Scaling the number of actors. Performance consistently improves as we scale the number of actors from 8 to 256, note that the number of learning updates performed does not depend on the number of actors.

Experiments

CONTINUOUS CONTROL

APE-X DPG

- 4 tasks
 - manipulator
 - humanoid : Standing / Walking / Running
- observation : features (small FC layers)
- 64 actors
 - ~14K total FPS
 - 86 batches of 256 transitions / sec (~22K transitions /sec)



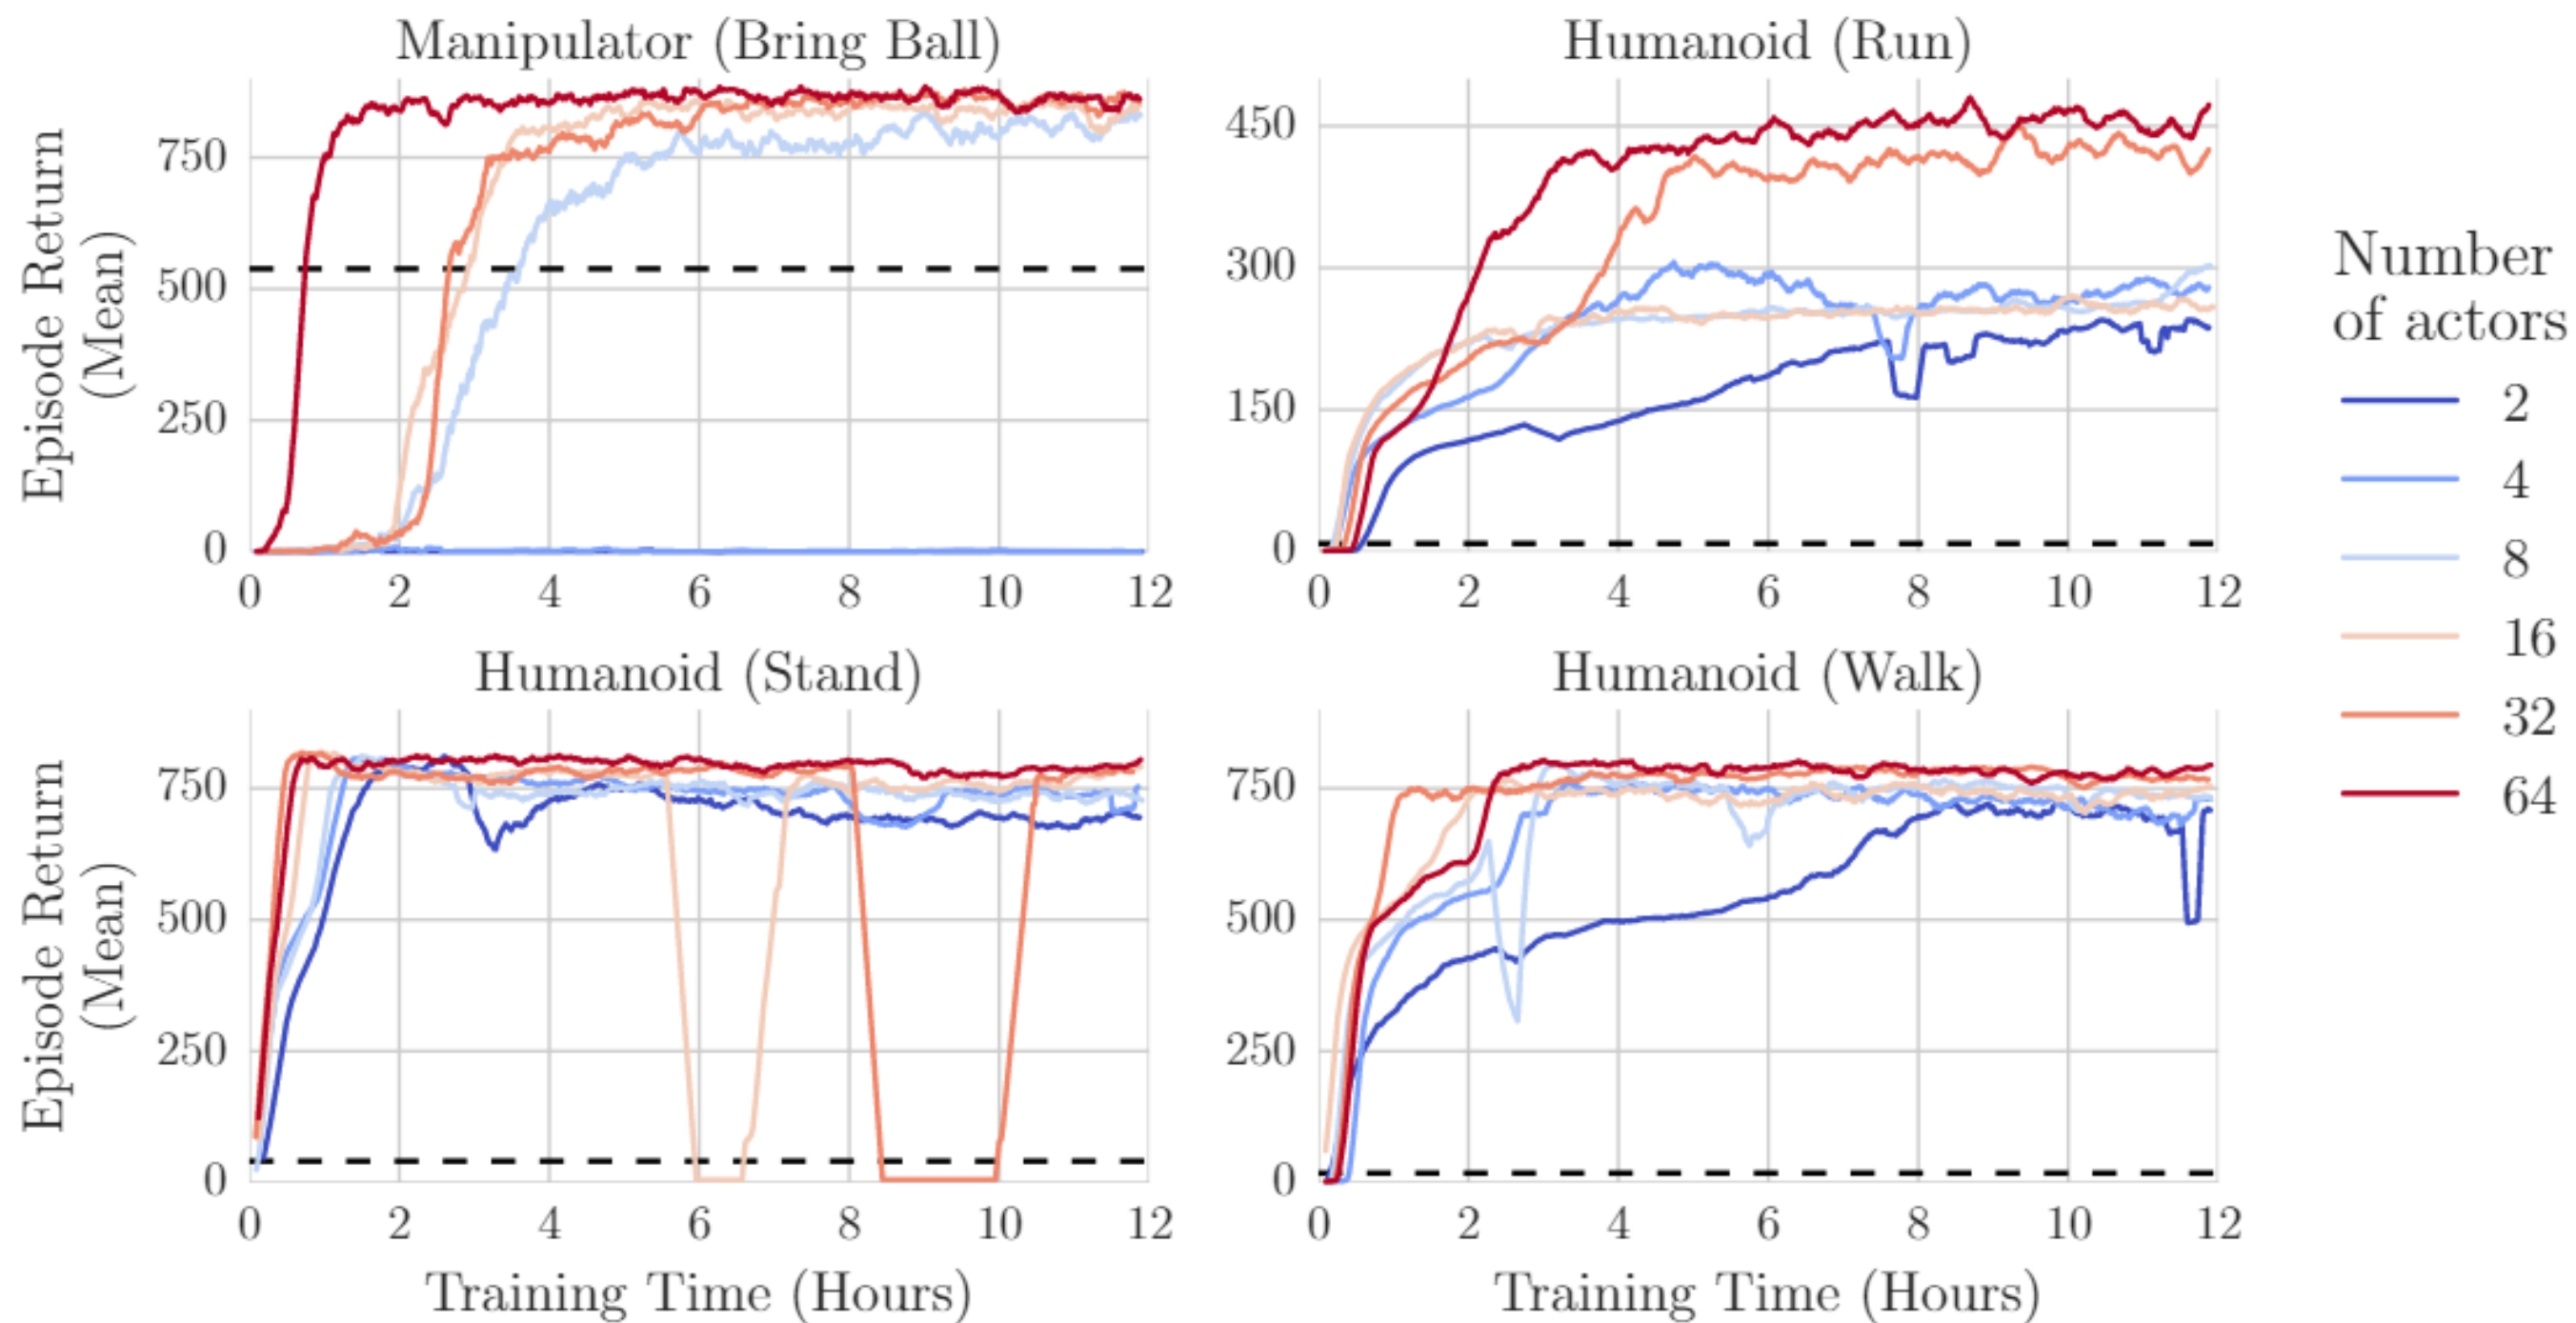


Figure 3: Performance of Ape-X DPG on four continuous control tasks, as a function of wall clock time. Performance improves as we increase the numbers of actors. The black dashed line indicates the maximum performance reached by a standard DDPG baseline over 5 days of training.

Analysis

HOW THE PERFORMANCE SCALES WITH THE NUMBER OF ACTORS

actor수가 많을수록 퍼포먼스가 좋을까?

- 8, 16, 32, 64, 128 and 256) for 35 hours on a subset of 6 Atari games
- shared experience replay memory 크기 고정 (1 million transition)
- 다다익선
- PER 없이도 비교해보았는데 actor가 많을수록 좋다는 결과는 동일했음(appendix)
- 본 연구의 아키텍처가 흔한 RL의 실패를 개선 시킬 수 있을 것
 - 어떤 실패?
 - the policy discovered is a **local optimum** in the parameter space, but **not a global one**
 - 왜 개선이 되는가?
 - large number of actors with varying amounts of exploration helps to **discover promising new courses of action**



Analysis

VARYING THE CAPACITY OF THE REPLAY MEMORY

replay memory의 용량이 달라지면 어떨까?

- 256 actors
- a median of ~37K total environment frames per second (약 ~9K transition)
- the contents of the memory is replaced much **faster** than in most DQN-like agents
- a **small benefit** to using a larger replay capacity
 - 왜? the value of keeping some high priority experiences around for longer and replaying them



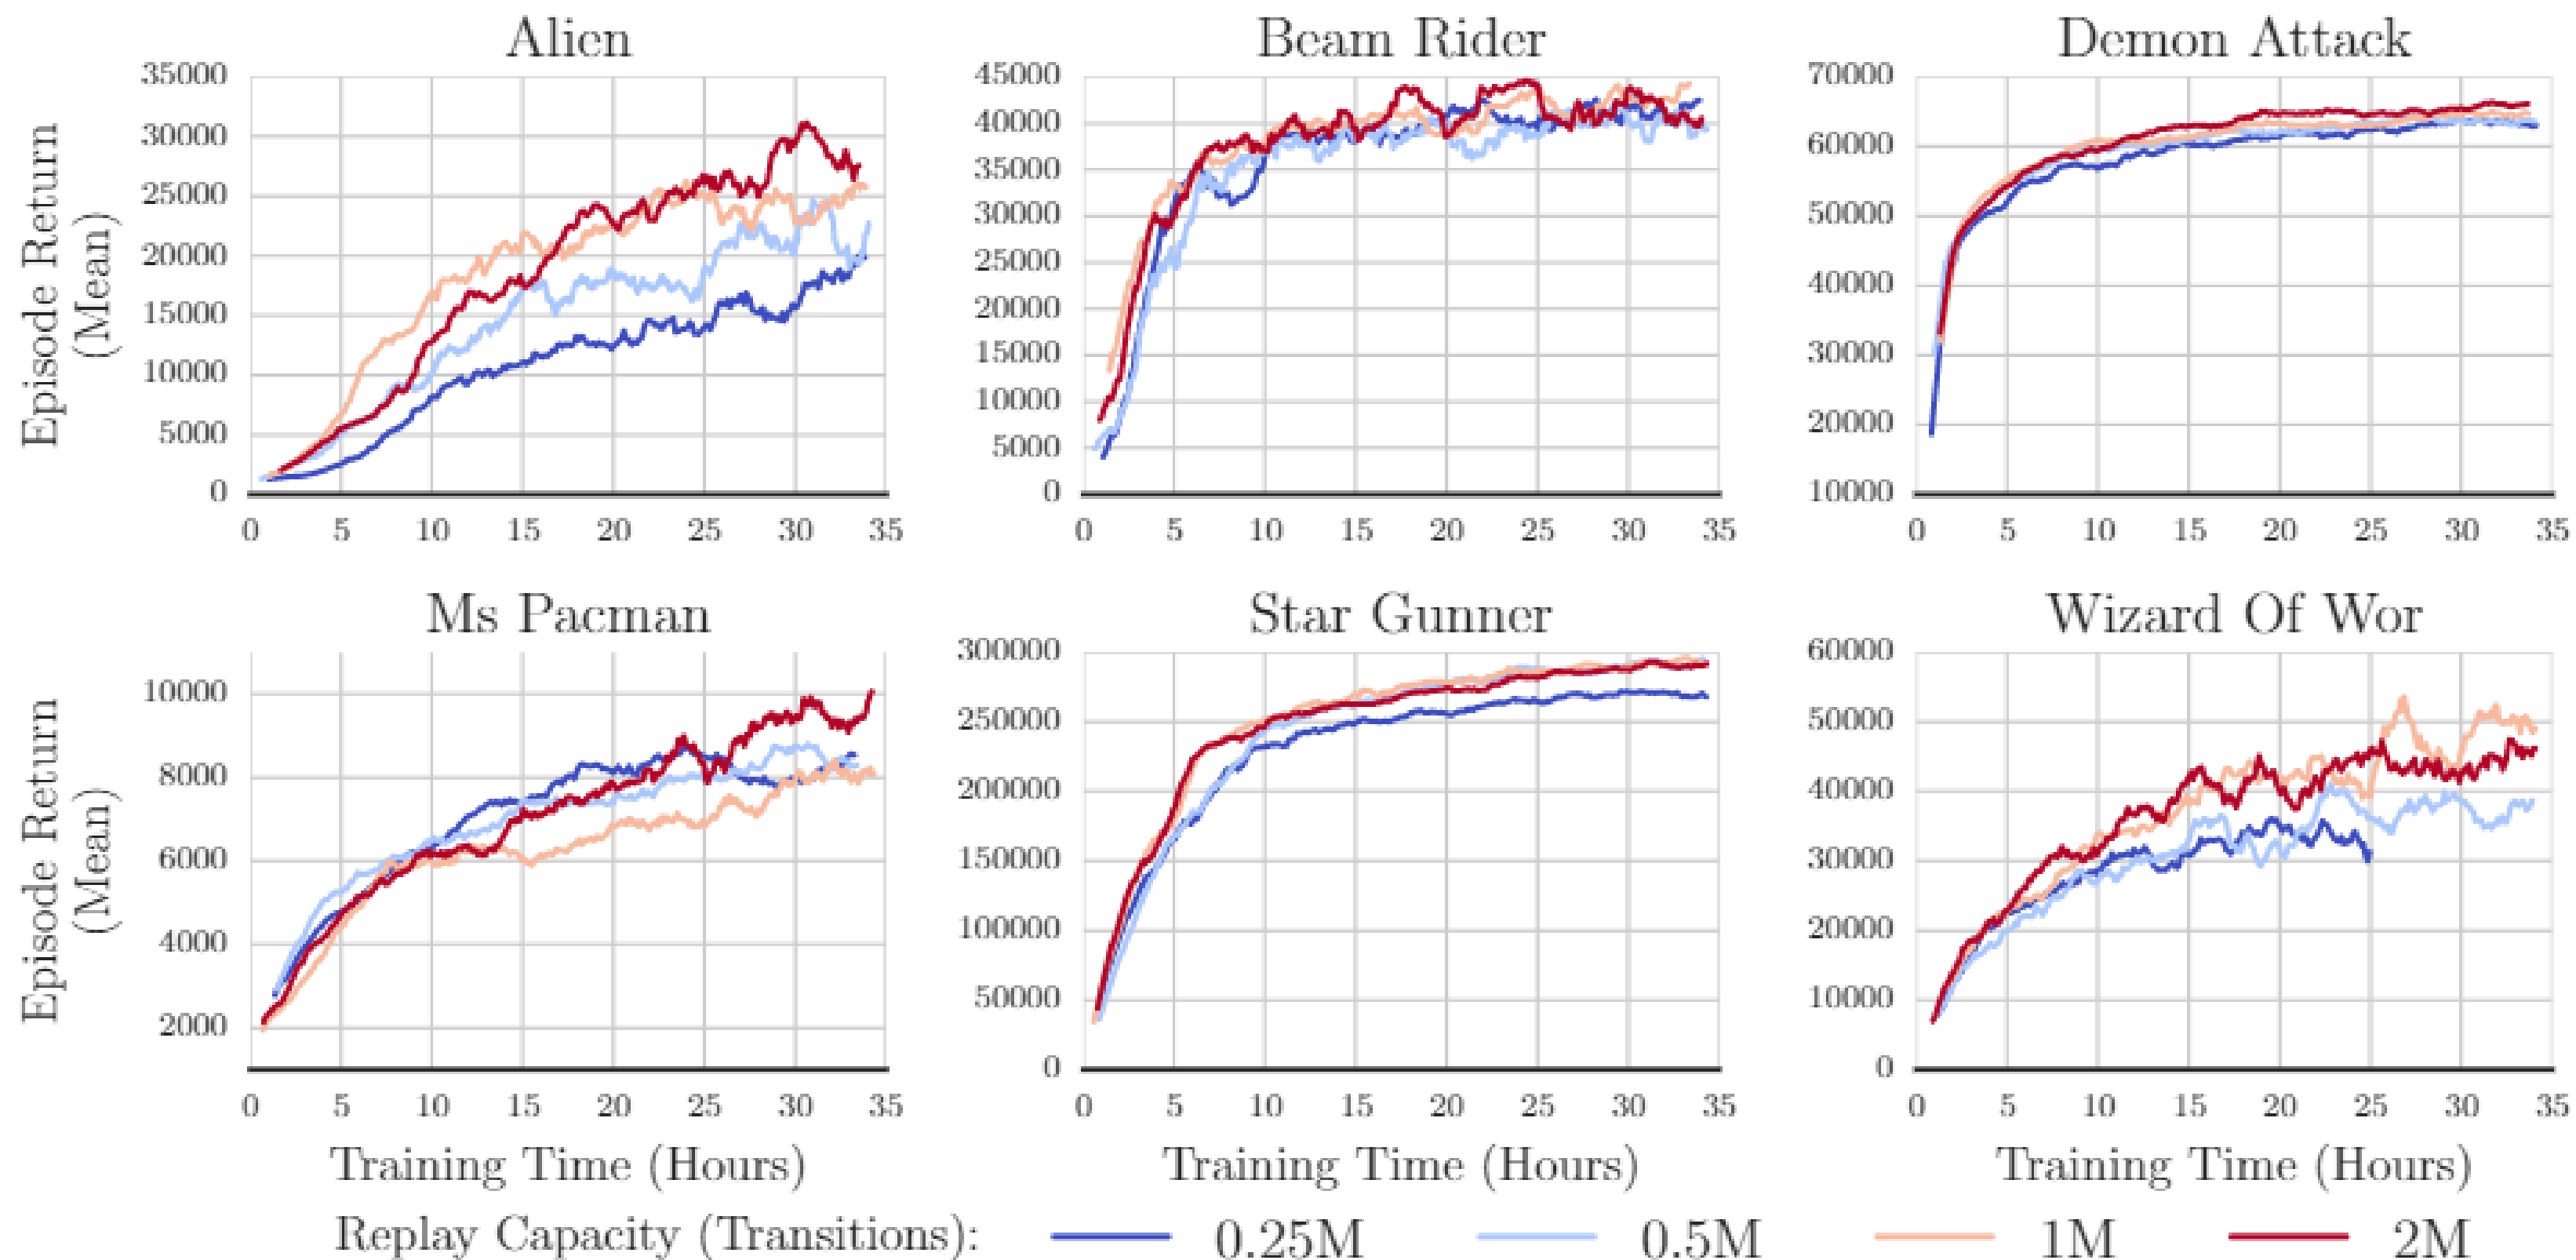


Figure 5: Varying the capacity of the replay. Agents with larger replay memories perform better on most games. Each curve corresponds to a single run, smoothed over 20 points. The curve for Wizard Of Wor with replay size 250K is incomplete because training diverged; we did not observe this with the other replay sizes.

Analysis

DISENTANGLE POTENTIAL EFFECTS OF TWO CONFOUNDING FACTORS IN OUR **SCALABILITY ANALYSIS**

Scalability의 다른 요인들

1. Recency of the experience data in the replay memory
2. Diversity of the data-generating policies

- appendix
- **neither** factor alone is sufficient to explain the performance we see
- gathering more experience data에 도움을 줄 뿐
- better exploration of the environment and better avoidance of overfitting



Conclusion

- Combined with **any other off-policy** reinforcement learning update
- 병렬적으로 많은 양의 데이터들을 만들어 낼 수 있는 상황에 적합한 체계 (many instances)
- RL의 문제인 large domain에서의 exploration 문제를
 - generating a diverse set of experiences = many actors
 - identifying and learning from the most useful events = per
- simple and direct approaches



Appendix




Figure 6: Testing whether improved performance is caused by recency alone: n denotes the number of actors, k the number of times each transition is replicated in the replay. The data in the run with $n = 32, k = 8$ is therefore as recent as the data in the run with $n = 256, k = 1$, but performance is not as good.

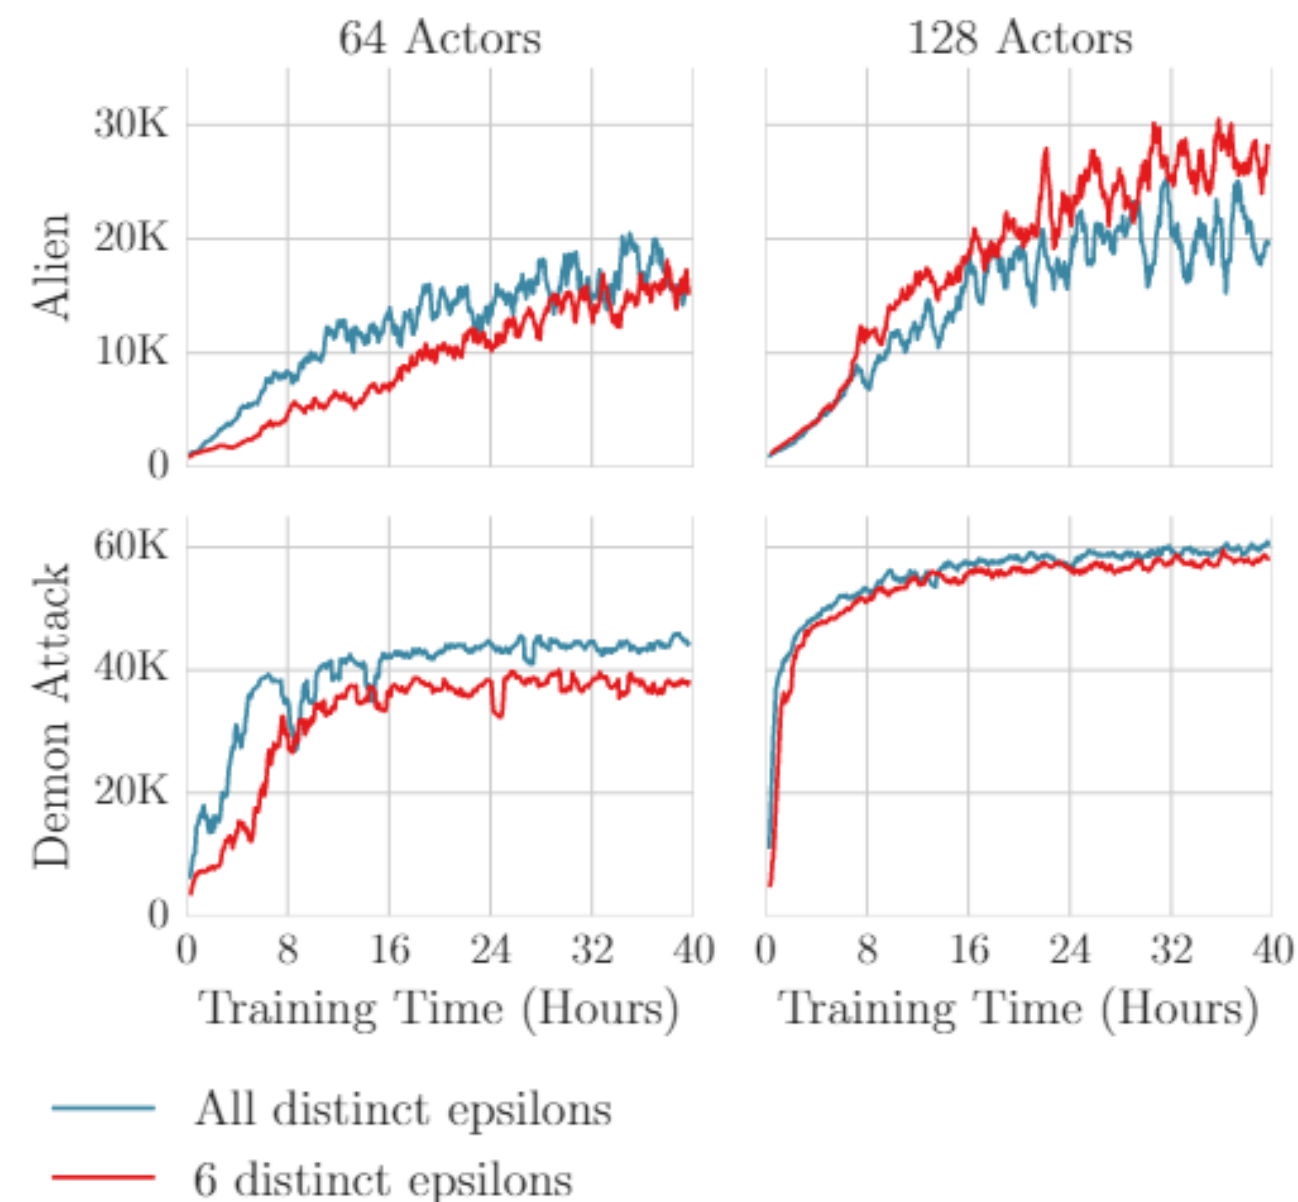


Figure 7: Varying the data-generating policies: **Red:** fixed set of 6 values for ϵ . **Blue:** full range of values for ϵ . In both cases, the curve plotted is from a separate actor that does not add data to the replay memory, and which follows an ϵ -greedy policy with $\epsilon = 0.00164$.

<https://www.slideshare.net/ssuser163469/distributed-prioritized-experience-replay>

<https://seolhokim.github.io/deeplearning/2020/04/09/apex-review/>

<https://www.opentutorials.org/module/3653/22071>