

Learning a Contact-Adaptive Controller for Robust, Efficient Legged Locomotion

Hansol Kang, Robotics Innovatory, SKKU




Contents

- ◆ Introduction
- ◆ Paper Review
- ◆ Implementation

Introduction

- ▶ Robots controlled through RL



► Learning a Contact-Adaptive Controller for Robust, Efficient Legged Locomotion

Learning a Contact-Adaptive Controller for Robust, Efficient Legged Locomotion

Xingye Da*, Zhaoming Xie*¹, David Hoeller*, Byron Boots*^{2,3},
Anima Anandkumar*⁴, Yuke Zhu*⁵, Buck Babich*, Animesh Garg**

Abstract: We present a hierarchical framework that combines model-based control and reinforcement learning (RL) to synthesize robust controllers for a quadruped (the Unitree Laikago). The system consists of a high-level controller that learns to choose from a set of primitives in response to changes in the environment and a low-level controller that utilizes an established control method to robustly execute the primitives. Our framework learns a controller that can adapt to challenging environmental changes on the fly, including novel scenarios not seen during training. The learned controller is up to 85 percent more energy efficient and is more robust compared to baseline methods. We also deploy the controller on a physical robot without any randomization or adaptation scheme.

Keywords: Legged Locomotion, Hierarchical Control, Reinforcement Learning

1 Introduction

Quadruped locomotion is often characterized in terms of *gaits* (walking, trotting, galloping, bounding, etc.) that have been well-studied in animals [1] and reproduced on robots [2, 3]. A gait is a periodic contact sequence that defines a specific contact timing for each foot. Controllers designed for these gaits have demonstrated robust behaviors on flat ground and rough terrain locomotion. However, it is rare to find controllers that can change gaits or contact sequences to adapt to environmental changes. An adaptive gait can reduce energy usage by removing unnecessary movement, as suggested in horse studies [1]. It is also required for completing more challenging scenarios such as riding a skateboard or recovery from leg slipping, as shown in Figure 1 (a, b).

In most model-based and learning-based control designs, the contact sequence is fixed or predefined [2, 3, 4, 5, 6, 7, 8]. Dynamic adaptation of the contact sequence is challenging because of the hybrid nature of legged locomotion dynamics as well as the inherent instability of such systems. While it is possible to generate adaptive contact schemes via trajectory optimization [9, 10, 11], such approaches are generally too compute-intensive for real-time use.

Here we present a hierarchical control framework for quadrupedal locomotion that learns to adaptively change contact sequences in real-time. A high-level controller is trained with reinforcement learning (RL) to specify the contact configuration of the feet, which is then taken as input by a low-level controller to generate ground reaction forces via quadratic programming (QP). At inference time, the high-level controller needs only evaluate a small multi-layer neural network, avoiding the use of an expensive model predictive control (MPC) strategy that might otherwise be required to optimize for long-term performance. The low-level controller provides high-bandwidth feedback to track base and foot positions and also helps ensure that learning is sample-efficient. The framework produces a controller that is up to 85 percent more energy efficient and also more robust than baseline approaches.

We train our controller with a simulated Unitree Laikago [12] on a split-belt treadmill, as shown in Figure 1 (c). The two belts can adjust speed independently, and we change the robot orientation to increase variation. In addition to comparing energy use and robustness to the baselines, we also demonstrate zero-shot transferability by testing the controller in novel situations such as one where a foot encounters a slippery surface (e.g., with zero friction), which we call the “banana peel” test.

*NVIDIA, ¹Univ. of British Columbia, ²Univ. of Washington, ³Caltech, ⁴UT Austin, ⁵Univ. of Toronto, Vector Institute. Work done at NVIDIA. Correspondence: xda@nvidia.com, zxie47@cs.ubc.ca

► 2020 Conference on Robot Learning (CoRL)

► 2020.11. arXiv


► Xingye Da¹, Zhaoming Xie^{1,2}, David Hoeller¹, Byron Boots^{1,3}, Anima Anandkumar^{1,4}, Yuke Zhu^{1,5}, Buck Babich¹, Animesh Garg^{1,6}

► ¹NVIDIA, ²Univ. of British Columbia, ³Univ. of Washington, ⁴Caltech, ⁵UT Austin, ⁶Univ. of Toronto


► Introduce a hierarchical control structure that combines model-based control design and model-free reinforcement learning for legged locomotion.

Paper Review

► How to control a robot?



1 DoF* link



1 DoF link

* Degree of Freedom

Paper Review

► Quadratic Programming(QP)

QP is an optimization tool

$$\text{minimize } \frac{1}{2} \mathbf{x}^T \mathbf{Q} \mathbf{x} + \mathbf{c}^T \mathbf{x}$$

$$\text{subject to } \mathbf{A} \mathbf{x} \leq \mathbf{b}$$

$$\text{minimize } 2x_1^2 + x_1x_2 + x_2^2 + x_1 + x_2$$

$$\text{subject to } x_1 + x_2 = 1$$

$$0 \leq x_1 \leq 0.7$$

$$0 \leq x_2 \leq 0.7$$

$$\text{minimize } \frac{1}{2} x^T \begin{bmatrix} 4 & 1 \\ 1 & 2 \end{bmatrix} x + \begin{bmatrix} 1 \\ 1 \end{bmatrix}^T x$$

$$\text{subject to } \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \leq \begin{bmatrix} 1 & 1 \\ 1 & 0 \\ 0 & 1 \end{bmatrix} x \leq \begin{bmatrix} 1 \\ 0.7 \\ 0.7 \end{bmatrix}$$

Paper Review

- How to use QP for quadruped robot control?

$$\underbrace{\begin{bmatrix} I & \cdots & I \\ [p_{com,1\times}] & \cdots & [p_{com,c\times}] \end{bmatrix}}_A \underbrace{\begin{bmatrix} f_1 \\ \vdots \\ f_c \end{bmatrix}}_f = \underbrace{\begin{bmatrix} m(\ddot{x}_{com}^d + g) \\ I_g \dot{\omega}_b^d \end{bmatrix}}_b$$

$$f^d = \arg \min_{f \in \mathbb{R}^k} (Af - b)^T S (Af - b) + \alpha f^T W f$$

$$s.t. \underline{d} < Cf < \bar{d}$$

$S \in \mathbb{R}^{6 \times 6}$ and $W \in \mathbb{R}^{k \times k}$ are positive – definite weight matrices

$$\underline{d}, \bar{d} \in \mathbb{R}^p, C \in \mathbb{R}^{p \times k}$$

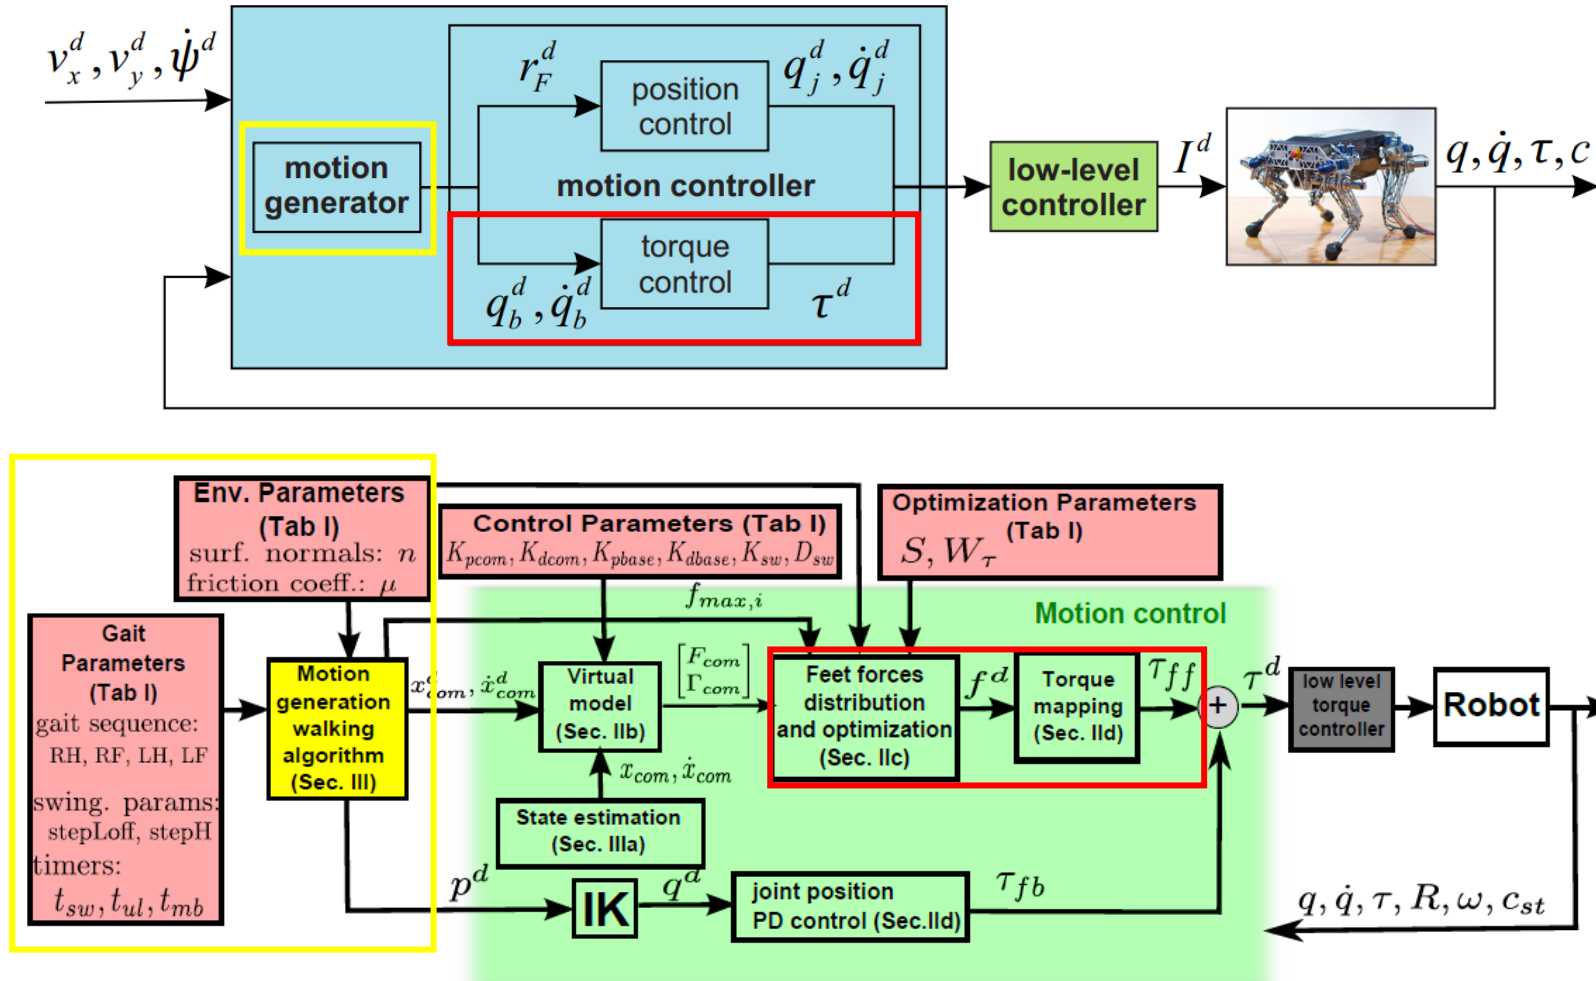
$k = 3 \times \text{Num. of contact}, p = \text{Num. of inequality const.}$

$\alpha \in \mathbb{R}$ weighs the secondary objective.



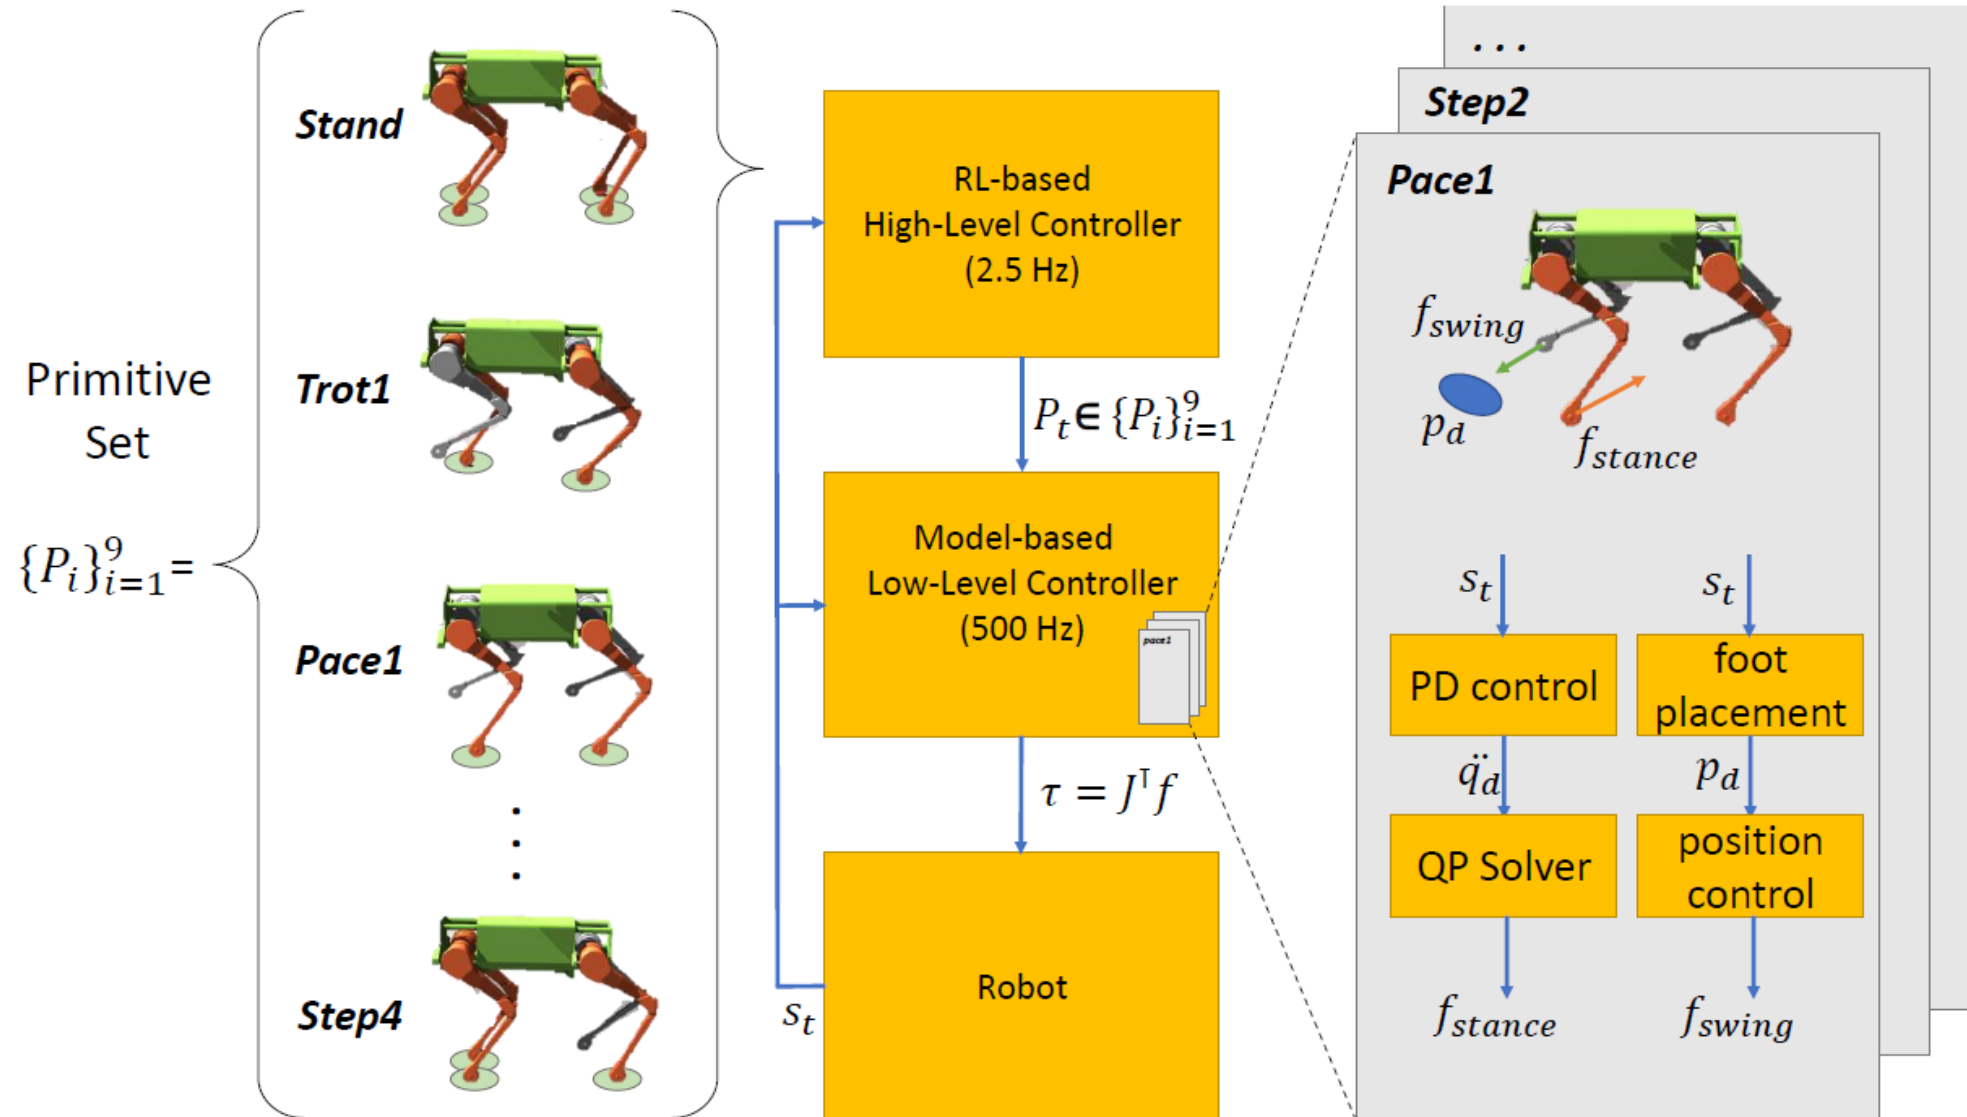
Paper Review

► How to use QP for quadruped robot control?



Paper Review

▶ Controller overview



Paper Review

► Primitive set



Stand



Trot 1



Trot 2



Pace 1



Pace 2




Step 1



Step 2



Step 3



Step 4

Paper Review

► High level controller

Simulator : Issac Gym

State Space

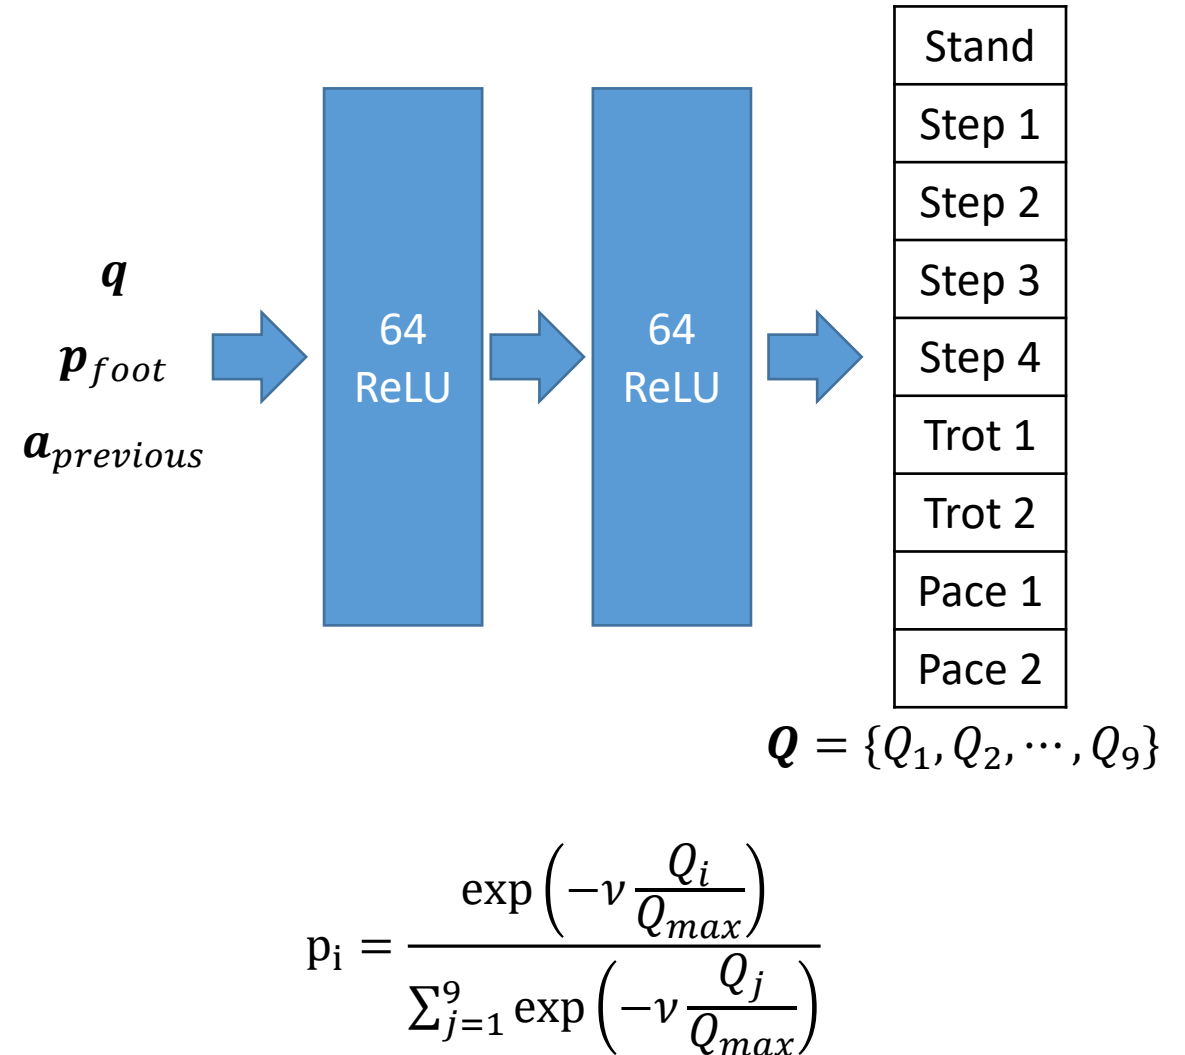
- Body pose \mathbf{q} (exclude x, y linear position)
- Relative foot position \mathbf{p}_{foot}
- Previously-used primitive $\mathbf{a}_{previous}$

Action Space

- 9 Primitives

Reward

$$r = 1 - 0.0025 \frac{1}{T} \sum \|\tau\|^2 - \frac{1}{T} \sum \|\dot{p}_{d,body} - \dot{p}_{body}\|^2$$



Paper Review

► High level controller

B Q-Learning Algorithm

We use DQN like algorithm to train our high-level policy. Details are shown in Algorithm 1.

Algorithm 1: Q Learning

initialization Q-function parameters θ_1, θ_2 for $Q_{\theta_1}, Q_{\theta_2}$, empty replay buffer D ;

set target network parameters $\theta_{\text{targ},1}, \theta_{\text{targ},2} \leftarrow \theta_1, \theta_2$ for $Q_{\theta_{\text{targ},1}}, Q_{\theta_{\text{targ},2}}$;

while *not done* **do**

 observe current state s ;

 sample action a based on Q-function;

 observe next state s' , reward r and done signal d ;

 store (s, a, r, d, s') in replay buffer D ;

if d is *True* or *time limit reached* **then**

 reset environment;

end

if *time to update* **then**

for $j = 1, 2, \dots$ *number of update* **do**

 sample batch of transition data $B = \{s, a, r, d, s'\}$;

 compute $a' = \arg \max_a Q_{\theta}(s', a)$;

 compute target $q_{\text{targ}} = r + (1 - d)\gamma \min_{i=1,2} (Q_{\theta_{\text{targ},i}}(s', a'))$;

 update θ_1, θ_2 by taking gradient descent w.r.t the objective function

$\frac{1}{|B|} \sum_{(s,a,r,d,s') \in B} ((Q_{\theta_1}(s, a) - q_{\text{targ}})^2 + (Q_{\theta_2}(s, a) - q_{\text{targ}})^2)$;

if $j \bmod 2 = 1$ **then**

$\theta_{\text{targ},1} \leftarrow \rho \theta_{\text{targ},1} + (1 - \rho) \theta_1$;

$\theta_{\text{targ},2} \leftarrow \rho \theta_{\text{targ},2} + (1 - \rho) \theta_2$;

end

end

end

end

Paper Review

► Low level controller

Base Pose Control

$$\min_f ||\mathbf{M}\mathbf{f} - \tilde{\mathbf{g}} - \ddot{\mathbf{q}}_d||\mathbf{Q} + ||\mathbf{f}||\mathbf{R}$$

subject to $f_{z,i} \geq f_{z,min}$ if $P_{t,i}$ is Stance

$f_{z,i} = 0$ if $P_{t,i}$ is Swing

$$-\mu f_x \leq f_z \leq \mu f_x$$

$$-\mu f_y \leq f_z \leq \mu f_y$$

Swing Foot Control

$$p_{d,i} = p_{0,i} + k(\dot{p}_{body} - \dot{p}_{d,body})$$


$$f_i = k_{p,i}(p_{d,i} - p_i) - k_{d,i}\dot{p}_i$$

Paper Review

► Result

Treadmill Speed 0 m/s

Step4
Step3
Step2
Step1
Pace2
Pace1
Trot2
Trot1
Stand




The controller chooses primitive "Stand" at zero speed

Paper Review

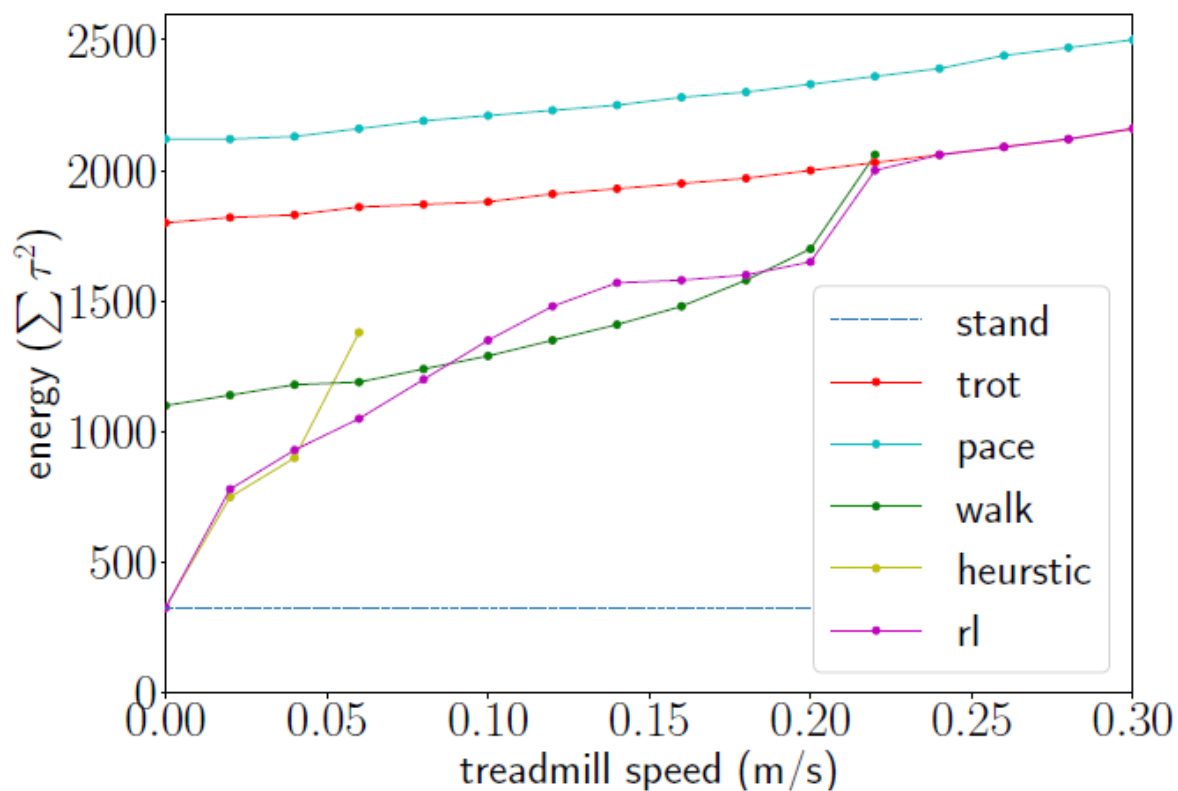
► Result

We show a controller that learns to adaptively change contact sequences in real time




Paper Review

Result




(a) Energy comparison over different speeds.



(b) Energy comparison over different yaws.

Paper Review

► Zero-shot adaptation



Paper Review

► Zero-shot adaptation



Conference on
Robot Learning



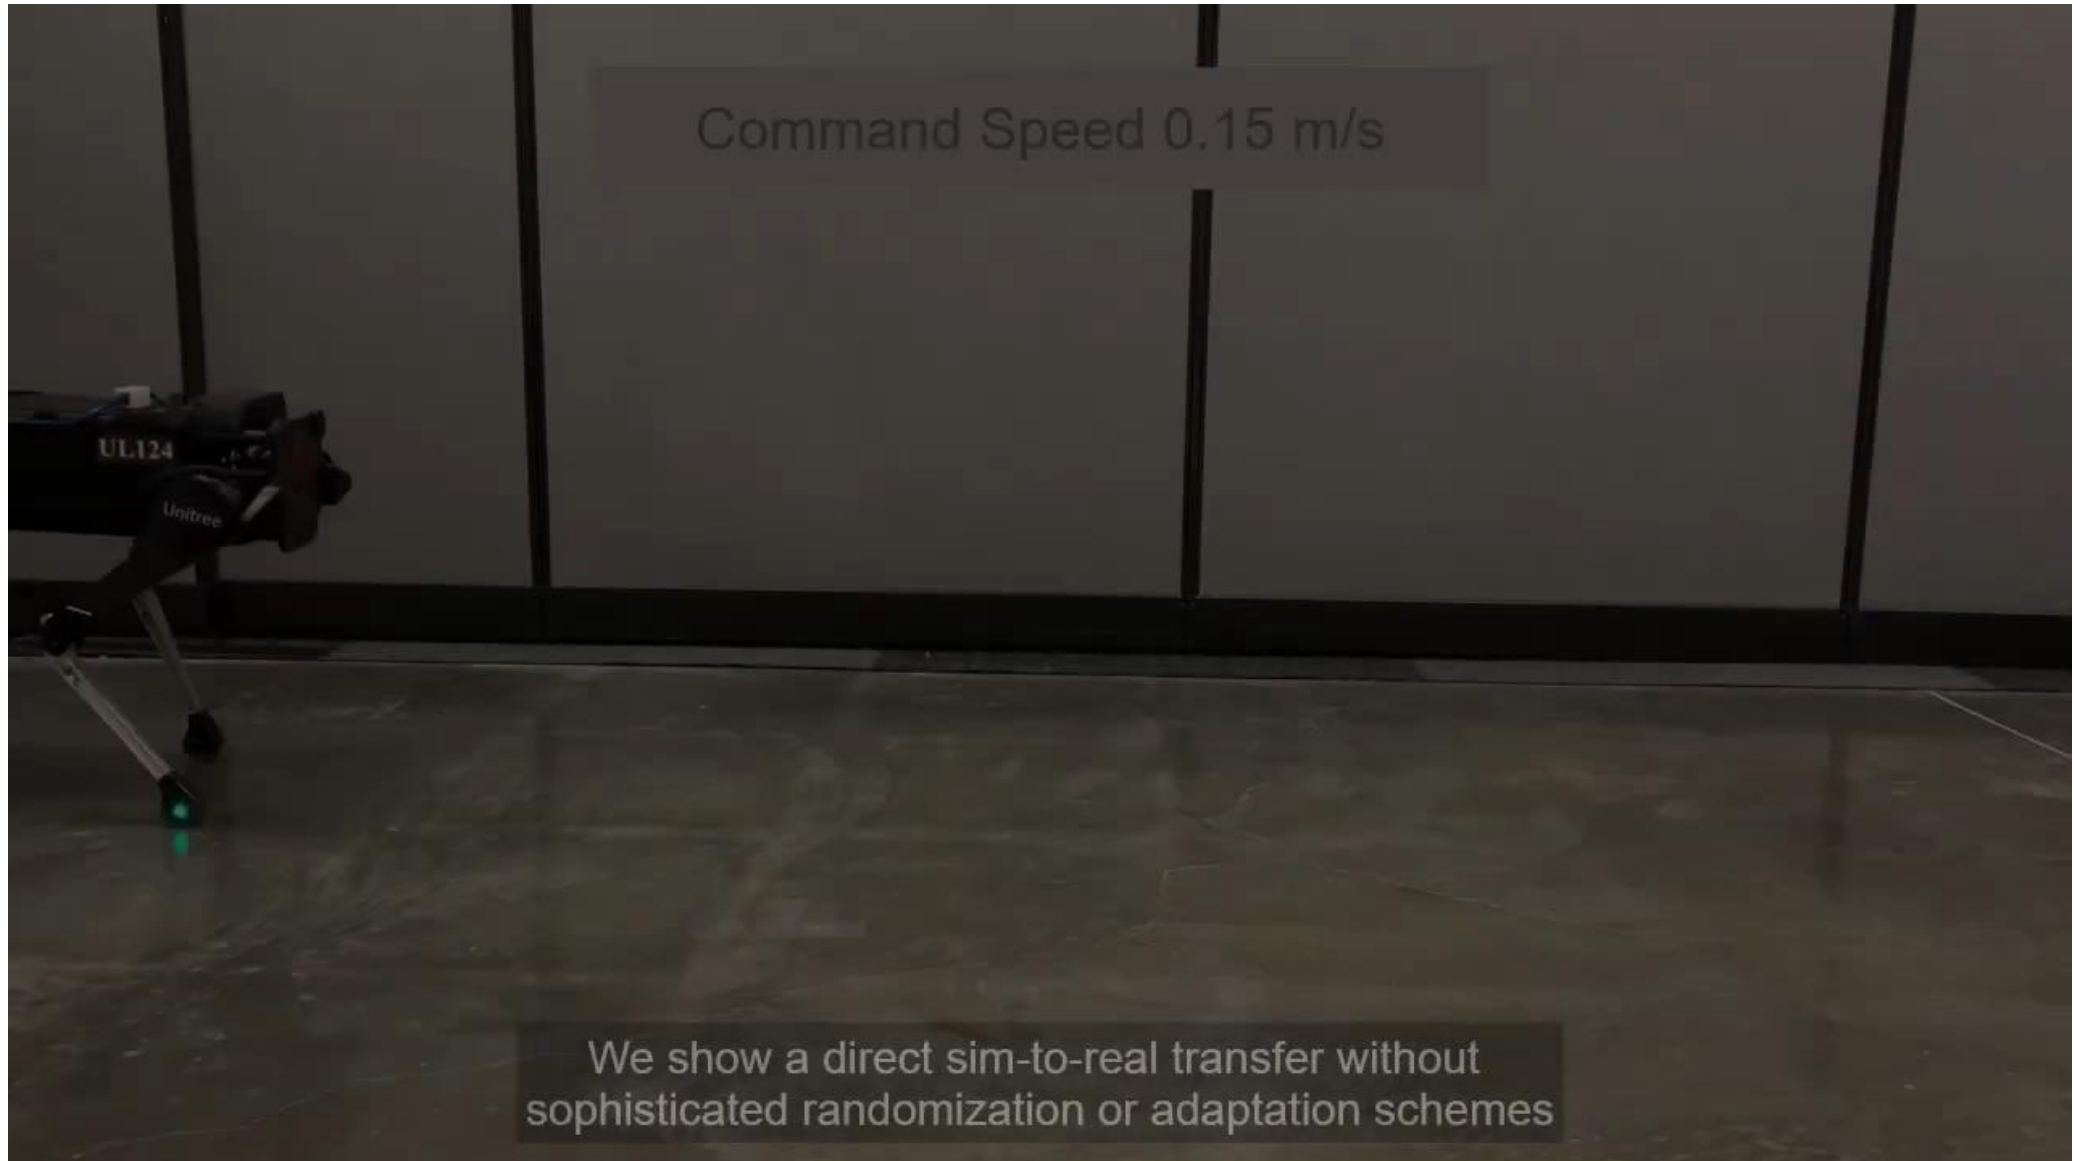
Extension: Skateboard Test



We show a preliminary result of Laikago skateboarding

Paper Review

► Sim-to-real test



Environment

The image shows a Linux desktop environment with a dark theme. The top panel displays system status: '월 18 : 5' (Monday 18:05), 'Ln 384, Col 46', 'Spaces: 4', 'UTF-8', 'LF', 'Python', and a battery icon. The main workspace is occupied by a Visual Studio Code window titled 'aidnvi_QP_3.py - 8.QP - Visual Studio Code'. The editor shows a file named 'onsoulkang@LAB: ~/aidn-bot/src/8.QP' with a large block of numerical data (likely a matrix or vector) and a 'target velocity : 0.3' line. Below the editor, the 'TERMINAL' pane shows the output of a program, including 'stopping threads', 'Thread with taskId 0 exiting', 'Thread TERMINATED', 'destroy semaphore', 'semaphore destroyed', '1 [[0.122689 0.02025', 'destroy main semaphore', 'main semaphore destroyed', 'finished', 'numActiveThreads = 0', 'btShutdownExampleBrowser stopping threads', 'Thread with taskId 0 exiting', 'Thread TERMINATED', 'destroy semaphore', 'semaphore destroyed', 'destroy main semaphore', 'main semaphore destroyed', and 'onsoulkang@LAB: ~/aidn-bot/src/8.QP\$'. To the right of the editor, a 'Bullet Physics ExampleBrowser using OpenGL3+ [btgl] Release build' window displays a 3D simulation of a four-legged robot on a checkered floor, with a shadow cast on the ground. The right sidebar of Visual Studio Code shows the 'EXPLORER' view with a file tree and the 'SOURCE CONTROL' view with a list of changes. The bottom status bar shows the current file path and encoding.

Q & A

Thank you for your attention