

RL² : FAST REINFORCEMENT LEARNING VIA SLOW REINFORCEMENT LEARNING

ICLR 2017

Yan Duan^{†‡}, John Schulman^{†‡}, Xi Chen^{†‡}, Peter L. Bartlett[†], Ilya Sutskever[‡], Pieter Abbeel^{†‡}

[†] UC Berkeley, Department of Electrical Engineering and Computer Science

[‡] OpenAI

TL;DR


Deep RL (DQN)

Score: **18.9**

experience measured in
real time: 40 days

“Slow”

vs.



Human

Score: 9.3

experience measured in
real time: **2 hours**


“Fast”

Why : The lack of good prior


Solution : RNN(GRU)

Environments

1. Multi-armed bandits



2. Tabular MDPs



3. ViZDoom



PRELIMINARIES(프리릴미너리스)

discrete-time finite-horizon discounted Markov decision process

$$\rho_{\mathcal{M}} : \mathcal{M} \rightarrow \mathbb{R}_+$$

$$\mathbf{M} = (S, A, P, r, \rho_0, \gamma, T)$$

bounded reward $r : S \times A \rightarrow [-R_{max}, R_{max}]$

↑
horizon

maximize $\eta(\pi_\theta) = E_\tau \left[\sum_{t=0}^T \gamma^t r(s_t, a_t) \right]$, where $\tau = (s_0, a_0, \dots)$

FORMULATION

learning an RL algorithm as a reinforcement learning \Rightarrow RL²

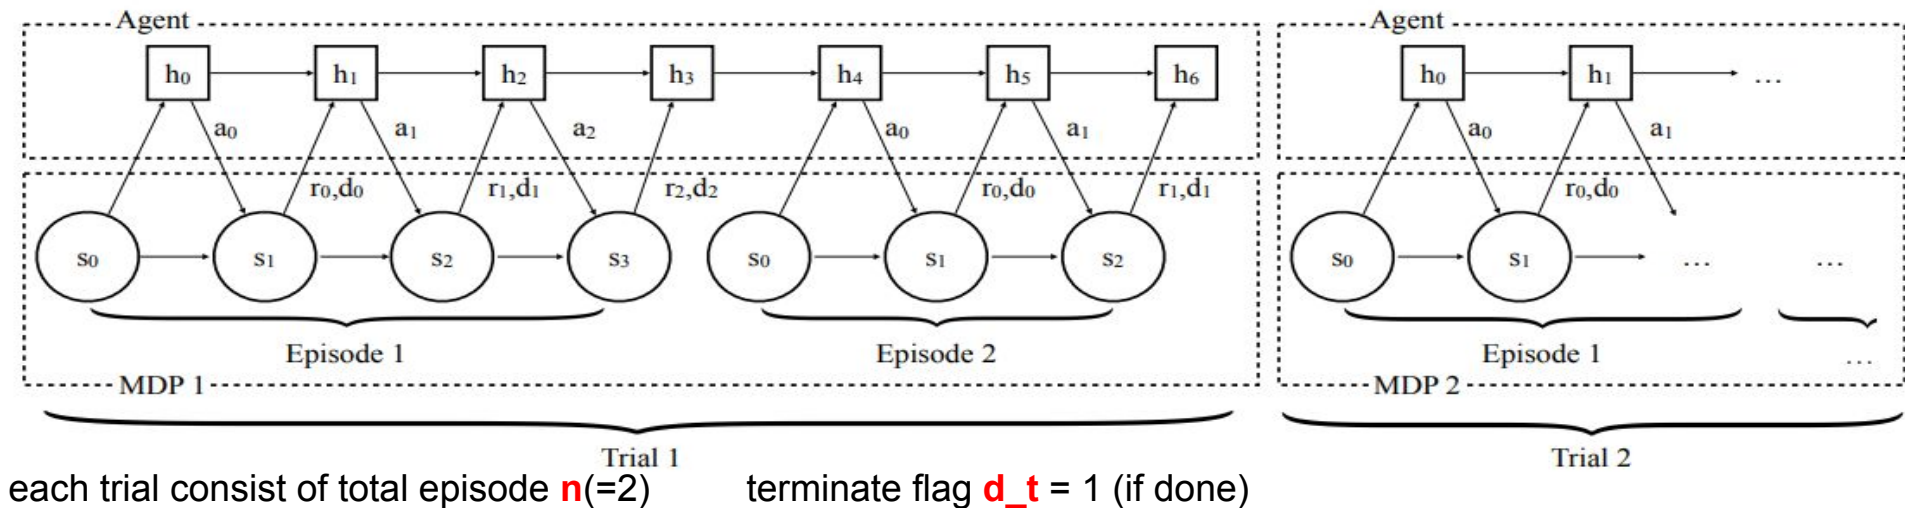


Figure 1: Procedure of agent-environment interaction

maximize the expected total discounted reward accumulated during a **single trial rather than a single episode**
minimizing the cumulative pseudo-regret 낮은 환경에서 빠른 적응력(=fast learning)을 획득

POLICY REPRESENTATION & OPTIMIZATION

embedding function $\phi(s, a, r, d)$



action distribution

To alleviate the difficulty of training RNNs due to vanishing and exploding gradients we use GRU

TRPO. Why? excellent empirical performance, does not require excessive hyperparameter tuning

To reduce variance => baseline(RNN using GRUs as building blocks), GAE(Generalized Advantage Estimation)

Evaluation

Q1 : RL²이 특정한 MDP에서 잘되었던 기존 알고리즘보다 좋을까?

Q2 : RL²이 high-dimensional task에서도 될까?

Q1 : RL^20이 특정한 MDP에서 잘되었던 기존 알고리즘보다 좋을까?

A : Multi Armed Bandit 와 tabular에서 잘된다면 인정


*MAB : subset of MDPs, agent's environment is stateless, **k** arms(actions), every time step, agent pulls one of the arms(say i), each arm is Bernoulli dist. with parameter p_i . The key challenge is E&E

We randomly generated MAB by sampling from uniform dist.


compared against :

Random, Gittins index, UCB1, Thompson sampling(TS, OTS), e-greedy, greedy


	Setup	Random	Gittins	TS	OTS	UCB1	ϵ -Greedy	Greedy	RL ²
n = episode	$n = 10, k = 5$	5.0	6.6	5.7	6.5	6.7	6.6	6.6	6.7
	$n = 10, k = 10$	5.0	6.6	5.5	6.2	6.7	6.6	6.6	6.7
	$n = 10, k = 50$	5.1	6.5	5.2	5.5	6.6	6.5	6.5	6.8
k = arms	$n = 100, k = 5$	49.9	78.3	74.7	77.9	78.0	75.4	74.8	78.7
	$n = 100, k = 10$	49.9	82.8	76.7	81.4	82.4	77.4	77.1	83.5
	$n = 100, k = 50$	49.8	85.2	64.5	67.7	84.3	78.3	78.0	84.9
total reward	$n = 500, k = 5$	249.8	405.8	402.0	406.7	405.8	388.2	380.6	401.6
	$n = 500, k = 10$	249.0	437.8	429.5	438.9	437.1	408.0	395.0	432.5
	$n = 500, k = 50$	249.6	463.7	427.2	437.6	457.6	413.6	402.8	438.9



(a) $n = 10$



(b) $n = 100$



(c) $n = 500$

Figure 2: RL² learning curves for multi-armed bandits. Performance is normalized such that Gittins index scores 1, and random policy scores 0.

Q1 : RL²이 특정한 MDP에서 잘되었던 기존 알고리즘보다 좋을까?

Tabular MDPs : MAB가 sequential decision making을 잘 나타내지 못하므로 사용.

compared against :

Random, (O)PSRL, BEB, UCRL2, e-greedy, greedy

$$|\mathcal{S}| = 10, |\mathcal{A}| = 5.$$


$T = 10$

Table 2: Random MDP Results

n = episode	Setup	Random	PSRL	OPSRL	UCRL2	BEB	ϵ -Greedy	Greedy	RL ²
$n = 10$		100.1	138.1	144.1	146.6	150.2	132.8	134.8	156.2
$n = 25$		250.2	408.8	425.2	424.1	427.8	377.3	368.8	445.7
$n = 50$		499.7	904.4	930.7	918.9	917.8	823.3	769.3	936.1
$n = 75$		749.9	1417.1	1449.2	1427.6	1422.6	1293.9	1172.9	1428.8
$n = 100$		999.4	1939.5	1973.9	1942.1	1935.1	1778.2	1578.5	1913.7

reward is Gaussian dist
with unit variance,
mean is sampled from
Normal(1,1)

Transitions are
sampled from flat
Dirichlet dist.



n이 높으면 잘안됨.

reinforcement learning
problem in the outer
loop becomes more
challenging to solve

Q2 : RL^20| high-dimensional task에서도 될까? + POMDP




Figure 4: Visual navigation. The target block is shown in red, and occupies an entire grid in the maze layout.

During multiple episodes, maze structure and target position are held fixed.

During each trial, we sample 1 out of 1000 randomly generated configurations of map layout and target positions

Reward : +1 reaches the target, -0.001 when hit the wall, -0.04 per time step. (Sparse reward)

The optimal strategy : explore the maze efficiently during the first episode,

and after locating the target, act optimally against the current maze and target based on the collected information

Table 3: Results for visual navigation. These metrics are computed using the best run among all runs shown in Figure 5. In 3c, we measure the proportion of mazes where the trajectory length in the second episode does not exceed the trajectory length in the first episode.

(a) Average length of successful trajectories



Episode	Small	Large
1	52.4 ± 1.3	180.1 ± 6.0
2	39.1 ± 0.9	151.8 ± 5.9
3	42.6 ± 1.0	169.3 ± 6.3
4	43.5 ± 1.1	162.3 ± 6.4
5	43.9 ± 1.1	169.3 ± 6.5

(b) %Success

Episode	Small	Large
1	99.3%	97.1%
2	99.6%	96.7%
3	99.7%	95.8%
4	99.4%	95.6%
5	99.6%	96.1%

(c) %Improved

Small	Large
91.7%	71.4%



each curve different random initialization of the RNN

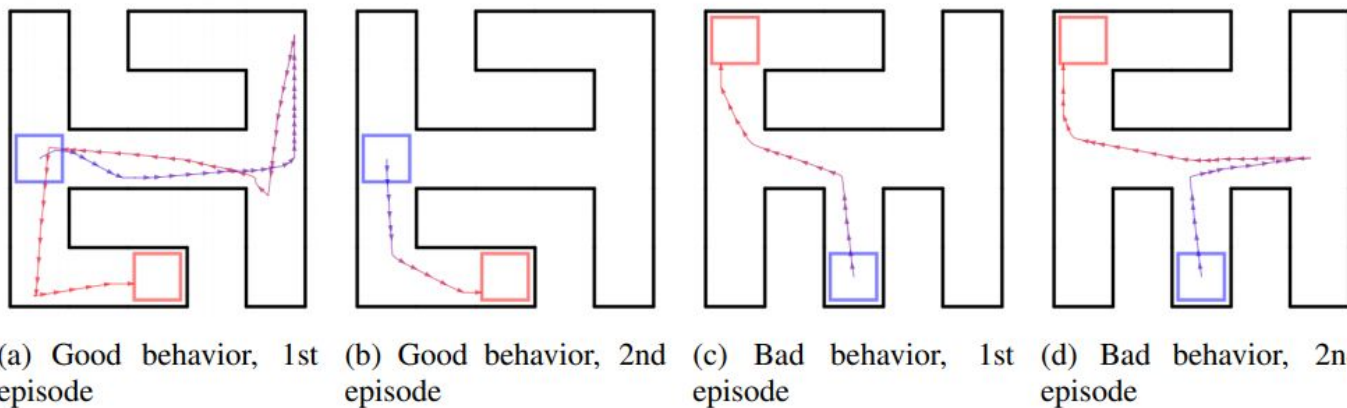


Figure 6: Visualization of the agent's behavior. In each scenario, the agent starts at the center of the blue block, and the goal is to reach anywhere in the red block.

a,b는 과거 정보를 잘 활용했으나 c,d는 과거정보를 활용못함.

논문에선 outer-loop에서 더좋은 RL techniques가 해결해줄거라고 믿음

Related work

1. Auto tuning of hyperparameters (Ishii et al., 2002)
2. Taylor & Stone (2009) survey the multi-task and transfer learning aspects
3. Fu et al. (2015) propose a model-based approach on top of iLQG with unknown dynamics (Levine & Abbeel, 2014), 이전 task에서 얻은 sample을 새로운 task에서 one-shot learning, but related tasks thanks to reduced sample complexity.
4. Deep neural networks for multi-task learning and transfer learning (Parisotto et al., 2015; Rusu et al., 2015; 2016a; Devin et al., 2016; Rusu et al., 2016b)
5. Our work draws inspiration from formulates meta-learning as an optimization problem, and can thus be optimized end-to-end via gradient descent(2016)