

BADGR: An Autonomous Self-Supervised Learning-Based Navigation System


Gregory Kahn, Pieter Abbeel, Sergey Levine
Berkeley AI Research (BAIR), University of California, Berkeley

Robotics and Automation Letters, 2021

Motivation

- Navigation for mobile robots is often regarded as primarily a **geometric problem**.

(Construct map \rightarrow Plan path on the map \rightarrow Track the planned path)



Motivation

- Navigation for mobile robots is often regarded as primarily a **geometric problem**.
(Construct map → Plan path on the map → Track the planned path)
- However, these methods **could not handle terrain properties and physical attributes** of the environment for deciding the traversability. (ex) tall grass)



Motivation

- Navigation for mobile robots is often regarded as primarily a **geometric problem**.
(Construct map → Plan path on the map → Track the planned path)
- However, these methods **could not handle terrain properties and physical attributes** of the environment for deciding the traversability. (ex) tall grass)

Let's learn the traversability from experience.

Contribution

The paper propose **BADGR**,

- an end-to-end **learning-based** mobile robot navigation system
- that can be trained entirely with **self- supervised**,
- **off-policy data** gathered in **real-world** environments,
- without any simulation or human supervision (**auto labeling**),
- and can **improve** as it gathers more data.
- **Model based RL** for decision making.
- Perform **point goal navigation** task in an **unknown** environment. **(no prior knowledge of the environment)**

BADGR: An Autonomous Self-Supervised Learning-Based Navigation System

Gregory Kahn, Pieter Abbeel, Sergey Levine



Method

- Mobile robot platform
- Data collection
- Self-supervised data labelling
- Predictive model
- Planning

Method : Mobile robot platform

- **Mobile robot platform**
- Data collection
- Self-supervised data labelling
- Predictive model
- Planning

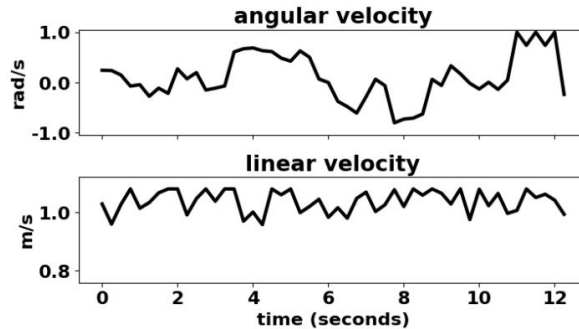
- Wheeled mobile robot
 - Sensor: Camera (for BADGR), 2D Lidar (for Baseline), IMU, GPS
 - Onboard computing: NVIDIA Jetson TX2
- linear acceleration,
angular velocity**



- Mobile robot platform
- **Data collection**
- Self-supervised data labelling
- Predictive model
- Planning

Method : Data collection

- Gather large amounts of diverse data with **minimal human intervention**.
- Used **time-correlated random walk** control policy for data gathering. (off-policy)
- Detect collision using Lidar and IMU. **If it is dangerous, reset the robot to safe coordinate.**



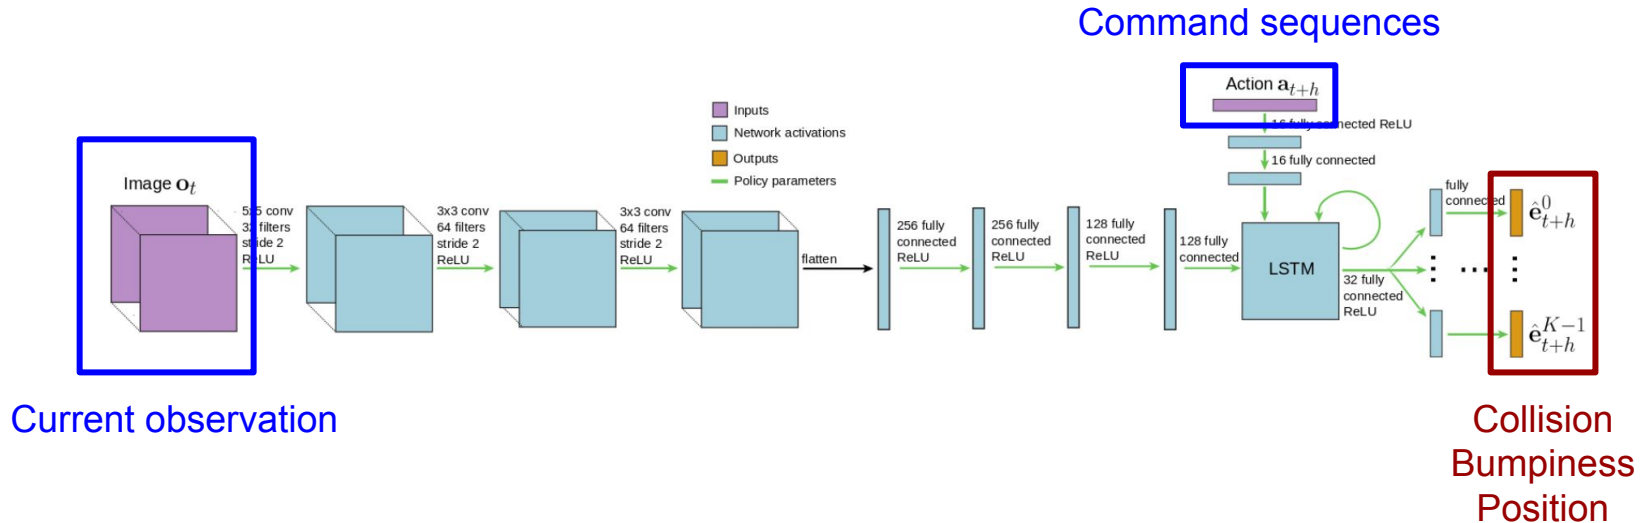
Method : Self-supervised data labelling

- Calculate labels for specific navigational events, which include **collision, bumpiness, position.**
- **Collision:** *Lidar < threshold || IMU magnitude sudden drop*
- **Bumpiness:** *|IMU angular velocity| > threshold*
- **Position:** *Wheel odometry + IMU*

- Mobile robot platform
- Data collection
- Self-supervised data labelling
- **Predictive model**
- Planning

Method : Predictive model

- **Predict H future steps**, conditioned on current observation and H intended future command sequences.



- Mobile robot platform
- Data collection
- Self-supervised data labelling
- **Predictive model**
- Planning

Method : Predictive model

- **Predict H future steps**, conditioned on current observation and H intended future command sequences.

Algorithm 1 Training BADGR

```

1: initialize dataset  $\mathcal{D} \leftarrow \emptyset$ 
2: while not done collecting data do
3:   get current observation  $\mathbf{o}_t$  from sensors
4:   get action  $\mathbf{a}_t$  from data collection policy
5:   add  $(\mathbf{o}_t, \mathbf{a}_t)$  to  $\mathcal{D}$ 
6:   execute  $\mathbf{a}_t$ 
7:   if in collision then
8:     execute reset maneuver
9:   end if
10: end while
11: for each  $(\mathbf{o}_t, \mathbf{a}_t) \in \mathcal{D}$  do
12:   calculate event labels  $\mathbf{e}_t^{0:K}$  using self-supervision
13:   add  $\mathbf{e}_t^{0:K}$  to  $\mathcal{D}$ 
14: end for
15: use  $\mathcal{D}$  to train predictive model  $f_\theta$  by minimizing Eqn. 1

```

Collect data

Self-labeling

Train

- Mobile robot platform
- Data collection
- Self-supervised data labelling
- Predictive model
- **Planning**

Method : Planning

- Gradient free optimizer + model predictive control
- Use **task specific rewards** which are functions of the outputs of learned predictive model

Prob. $\mathbf{a}_{t:t+H}^* = \arg \max_{\mathbf{a}_{t:t+H}} R(f_{\theta}(\mathbf{o}_t, \mathbf{a}_{t:t+H})),$

Sol.


$$\epsilon_h^n \sim \mathcal{N}(0, \sigma \cdot \mathbf{I}) \quad \forall n \in \{0 \dots N-1\}, h \in \{0 \dots H-1\}$$

$$\tilde{\mathbf{a}}_h^n = \beta \cdot \underbrace{(\hat{\mathbf{a}}_{h+1} + \epsilon_h^n)}_{\text{previous optimized distribution}} + \underbrace{(1 - \beta) \cdot \tilde{\mathbf{a}}_{h-1}^n}_{\text{time correlation}} \quad \text{where } \tilde{\mathbf{a}}_{h < 0} = 0$$

$$\tilde{R}^n = R(f_{\theta}(\mathbf{o}_t, \tilde{\mathbf{a}}_{0:H}^n))$$

$$\hat{\mathbf{a}}_{0:H} = \frac{\sum_{n=0}^N \exp(\gamma \cdot R^n) \cdot \tilde{\mathbf{a}}_{0:H}^n}{\sum_{n'=0}^N \exp(\gamma \cdot R^{n'})},$$

(Model based RL)



- Mobile robot platform
- Data collection
- Self-supervised data labelling
- Predictive model
- **Planning**

Method : Planning


- Gradient free optimizer + model predictive control
- Use **task specific rewards** which are functions of the outputs of learned predictive model

Prob. $\mathbf{a}_{t:t+H}^* = \arg \max_{\mathbf{a}_{t:t+H}} R(f_{\theta}(\mathbf{o}_t, \mathbf{a}_{t:t+H})),$

Sol. $\epsilon_h^n \sim \mathcal{N}(0, \sigma \cdot \mathbf{I}) \forall n \in \{0 \dots N-1\}, h \in \{0 \dots H-1\}$
 $\tilde{\mathbf{a}}_h^n = \beta \cdot (\underbrace{\hat{\mathbf{a}}_{h+1} + \epsilon_h^n}_{\text{previous optimized distribution}}) + (1 - \beta) \cdot \underbrace{\tilde{\mathbf{a}}_{h-1}^n}_{\text{time correlation}} \text{ where } \tilde{\mathbf{a}}_{h<0} = 0$

$$\tilde{R}^n = R(f_{\theta}(\mathbf{o}_t, \tilde{\mathbf{a}}_{0:H}^n))$$

$$\hat{\mathbf{a}}_{0:H} = \frac{\sum_{n=0}^N \exp(\gamma \cdot R^n) \cdot \tilde{\mathbf{a}}_{0:H}^n}{\sum_{n'=0}^N \exp(\gamma \cdot R^{n'})},$$



Algorithm 2 Deploying BADGR

- 1: **input:** trained predictive model f_{θ} , reward function R
- 2: **while** task is not complete **do**
- 3: get current observation \mathbf{o}_t from sensors
- 4: solve Eqn. 2 using f_{θ} , \mathbf{o}_t , and R
to get the planned action sequence $\mathbf{a}_{t:t+H}^*$
- 5: execute the first action \mathbf{a}_t^*
- 6: **end while**

Experiment

- Urban environment
- Off-road environment (e.g. tall grass)
- Self-improvement
- Generalization

Reward for “Point goal navigation”

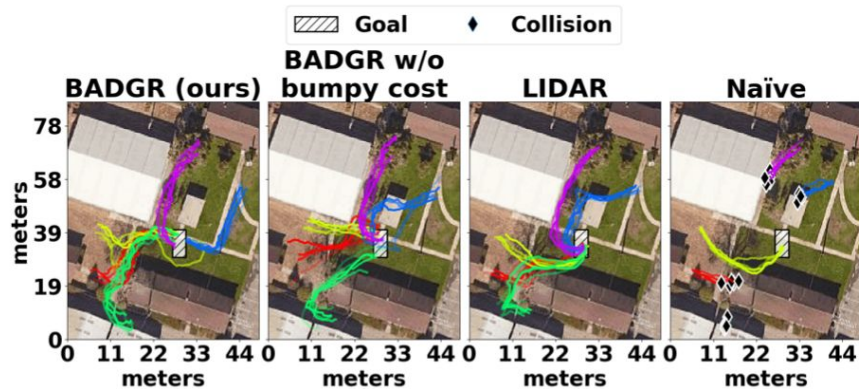
$$R^{\text{COLL}}(\hat{\mathbf{e}}_{t'}^{0:K}) = \hat{\mathbf{e}}_{t'}^{\text{COLL}} \quad \text{Don't collide}$$

$$R^{\text{POS}}(\hat{\mathbf{e}}_{t'}^{0:K}) = (1 - \hat{\mathbf{e}}_{t'}^{\text{coll}}) \cdot \frac{1}{\pi} \angle(\hat{\mathbf{e}}_{t'}^{\text{POS}}, \mathbf{p}^{\text{GOAL}}) + \hat{\mathbf{e}}_{t'}^{\text{coll}} \quad \text{Move to the goal direction}$$


$$R^{\text{BUM}}(\hat{\mathbf{e}}_{t'}^{0:K}) = (1 - \hat{\mathbf{e}}_{t'}^{\text{coll}}) \cdot \hat{\mathbf{e}}_{t'}^{\text{BUM}} + \hat{\mathbf{e}}_{t'}^{\text{coll}}, \quad \text{Don't run on bumpy terrain}$$

- **Urban environment**
- Off-road environment (e.g. tall grass)
- Self-improvement
- Generalization

Experiment : Urban environment


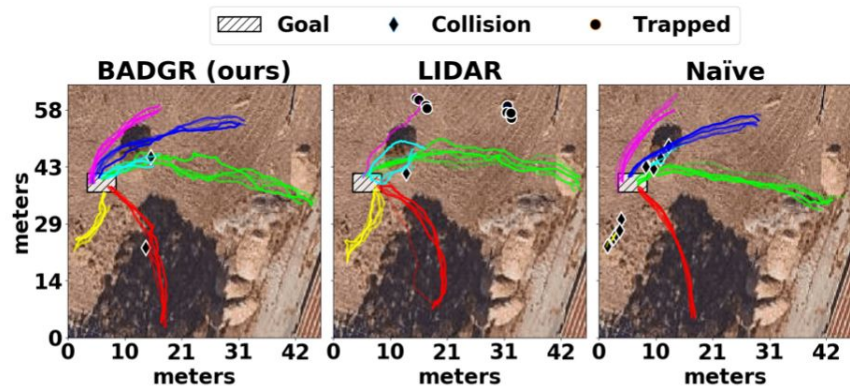


Method	Successfully Reached Goal	Avg. Bumpiness
BADGR (ours)	25/25 (100%)	8.7 ± 4.4
BADGR w/o bumpy cost	25/25 (100%)	15.0 ± 3.4
LIDAR	25/25 (100%)	13.3 ± 2.9
Naïve	5/25 (20%)	N/A




Experiment : Off-road environment

- Urban environment
- **Off-road environment (e.g. tall grass)**
- Self-improvement
- Generalization




Experiment : Self-improvement

- Urban environment
- Off-road environment (e.g. tall grass)
- **Self-improvement**
- Generalization



Experiment : Generalization

- Urban environment
- Off-road environment (e.g. tall grass)
- Self-improvement
- **Generalization**



Limitations (Personal thought)

- **Generalization** in various environments is left as question.
- **Directly** running robots **in real world** to collect data is inefficient and not good for robots.
- Difficult to be **robust** for various **goal positions**. (Typical problem for autonomous navigation in unknown environments)
- Wobbling control outputs requires **post-processing**.

**However, it is very simple
and works well in an unknown environment.**

Project page

<https://sites.google.com/view/badgr>

Q & A