

Neural Combinatorial Optimization with RL

백승언

08 Nov, 2021

- **Introduction**

- Combinatorial optimization problem
 - Examples of CO

- **Neural Combinatorial Optimization with Reinforcement Learning(NCO)**

- Backgrounds
- Pointer Network
- Methodology

- **Experiment results**

Introduction

Combinatorial optimization problems

● What is Combinatorial Optimization(CO) problem?

- Finding an optimal objects from a finite set of objects^[1]
- In many hard CO problems, exhaustive search is not tractable
 - Optimal solution require exponential computing time
- Conventional CO problems generally rely on handcrafted heuristics or optimization methods
 - Once the problem statement changes slightly, they need to be revised and this processes are time-consuming

● Machin learning approaches to CO problem

- Supervised learning is not applicable to most CO problems because it is difficult to acquire optimal labels
 - Exceptionally, Pointer Network beats the CO problem with small-size problems
 - TSP-50, TSP-100, TSP-200 and so on
- Many studies show that the Reinforcement Learning(RL) have potential to tackle the CO problems
 - This paper^[2] is the beginning of such studies.

[1] : https://en.wikipedia.org/wiki/Combinatorial_optimization

[2]: <https://arxiv.org/pdf/1611.09940.pdf>

Examples of CO

● Knapsack Problem(KP)

- 무게 제한 아래에서, 최고의 효용을 얻을 수 있는 물품들을 가방에 담는 문제


- $z = \sum_{i=1}^n v_i x_i$, s.t. $\sum_{i=1}^n w_i x_i \leq W$, $x_i \in \mathbb{Z}^{0+}$



● Travelling Salesman Problem(TSP)

- 방문 판매원이 최소 비용으로 모든 지역을 순회하는 방법에 관해 묻는 문제

- $z = ||x_n - x_1||_2 + \sum_{i=1}^{n-1} ||x_i - x_{i+1}||_2$




Neural Combinatorial Optimization with RL

- **The difficulty in applying existing search heuristics to newly encountered problem**
 - All search algorithms require the **prior knowledge** over problems to guarantee performance
 - Requirements for generalized solution at which can handle the various CO problems has increased
 - Oh, can neural network do this?
- **Previous application of neural networks to combinatorial optimization**
 - Hopfield network
 - Hopfield proposed Hopfield networks to solve the TSP (1985)
 - Limitation of Hopfield network was issued and resolved by Wilson & Pawley (1988)
 - Deformable template model
 - Durbin proposed the elastic network to solve TSP (1987)
 - The application of Self Organizing Map to TSP was proposed by Fort, Angeniol, Kohonen, (1988~1990)
 - Seq2Seq learning
 - Yuitan et al, Zoph & Le purposed the study for optimization in various domain. (2016)
 - Vinyals et al revisited the TSP by their **Pointer Network** (2015)

Important previous study: Pointer network

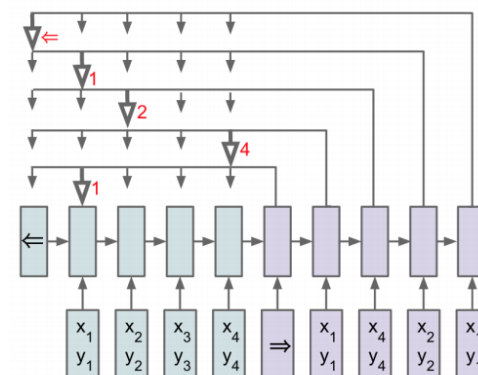
- A network that can adjust the **output length**

- Previous recurrent network cannot handle the variable frame
 - Seq2Seq paradigm beaten this issue using **encoder-decoder structure**
- Seq2Seq, however, has limitation that it cannot adjust the variable output length
 - The Pointer Network resolve this issue by simple idea
- The Pointer Network utilize the attention score of the input embedding vector
 - It seems like decoder output **point** the input vector



Sequence-to-Sequence / Content-based input attention

$$\begin{aligned}u_j^i &= v^T \tanh(W_1 e_j + W_2 d_i) \\a_j^i &= \text{softmax}(u_j^i) \\d_i' &= \sum_{j=1}^n a_j^i e_j \\p(\mathcal{C}^{\mathcal{P}} | \mathcal{P}; \theta) &= \prod_{i=1}^{m(\mathcal{P})} p_{\theta}(C_i | C_1, \dots, C_{i-1}, \mathcal{P}; \theta).\end{aligned}$$



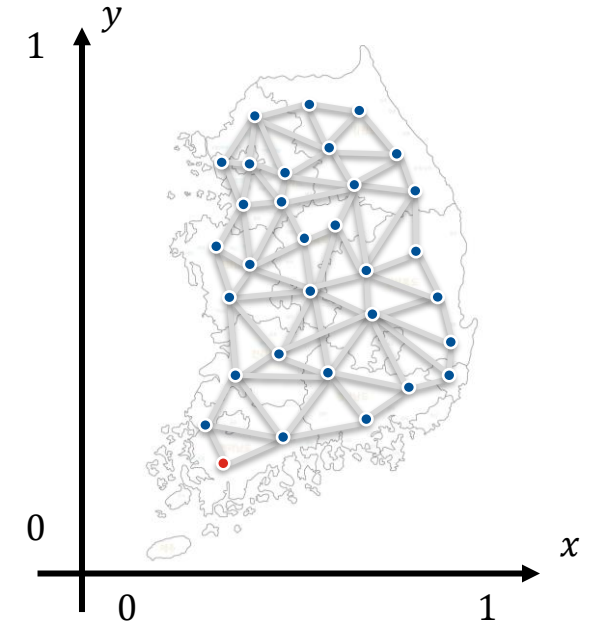
Pointer Network

$$\begin{aligned}u_j^i &= v^T \tanh(W_1 e_j + W_2 d_i) \\p(C_i | C_1, \dots, C_{i-1}, \mathcal{P}) &= \text{softmax}(u^i)\end{aligned}$$


● Neural network architecture for TSP

■ Problem definition

- State
 - sequence of n cities in a 2D Euclidean space
 - $s = \{\mathbf{x}_i\}_{i=1}^n, \mathbf{x}_i \in \mathbb{R}^2$
- Objective function of TSP(return G)
 - length of a tour defined by a permutation π
 - $L(\pi | s) = \|\mathbf{x}_{\pi(n)} - \mathbf{x}_{\pi(1)}\|_2 + \sum_{i=1}^{n-1} \|\mathbf{x}_{\pi(i)} - \mathbf{x}_{\pi(i+1)}\|_2, \|\cdot\|$ denotes l_2 norm
- Policy
 - Using Pointer network
 - Chain rule to factorize the probability of a tour as
 - $p(\pi | s) = \prod_{i=1}^n p(\pi(i) | \pi(< i), s)$
- Pointing mechanism
 - $u_i = \begin{cases} v^T \cdot \tanh(W_{ref} \cdot r_i + W_q \cdot q) & \text{if } i \neq \pi(j) \text{ for all } j < i \\ -\infty & \text{for } i = 1, 2, \dots, k \end{cases}$
 - $A(ref, q; W_{ref}, W_q, v) \triangleq \text{softmax}(C \tanh(\mathbf{u}))$, C is a hyperparameter



TSP on Euclidean space



Pointer network architecture

● Neural network architecture for TSP

- Policy optimization
 - Actor
 - Using Pointer network
 - Mapping input sequence \mathbf{s} and previous action set $\{a_j, j \in \{1, \dots, i-1\}\}$ to action a_i
 - Critic
 - Using Pointer network
 - Mapping input sequence \mathbf{s} to baseline $b_{\theta_v}(\mathbf{s})$
 - Loss function of actor, critic
 - $J(\theta | s) = \mathbb{E}_{\pi \sim p_\theta(\cdot | s)} L(\pi | s), \quad s \in \mathcal{S}$
 - $\nabla_\theta J(\theta | s) = \mathbb{E}_{\pi \sim p_\theta(\cdot | s)} [(L(\pi | s) - b(s)) \nabla_\theta \log p_\theta(\pi | s)],$
(REINFORCE)
 - $L(\theta_v) = \frac{1}{B} \sum_{i=1}^B \|b_{\theta_v}(s_i) - L(\pi_i | s)\|_2^2$

Algorithm 1 Actor-critic training

```
1: procedure TRAIN(training set  $S$ , number of training steps  $T$ , batch size  $B$ )
2:   Initialize pointer network params  $\theta$ 
3:   Initialize critic network params  $\theta_v$ 
4:   for  $t = 1$  to  $T$  do
5:      $s_i \sim \text{SAMPLEINPUT}(S)$  for  $i \in \{1, \dots, B\}$ 
6:      $\pi_i \sim \text{SAMPLESOLUTION}(p_\theta(\cdot | s_i))$  for  $i \in \{1, \dots, B\}$ 
7:      $b_i \leftarrow b_{\theta_v}(s_i)$  for  $i \in \{1, \dots, B\}$ 
8:      $g_\theta \leftarrow \frac{1}{B} \sum_{i=1}^B (L(\pi_i | s_i) - b_i) \nabla_\theta \log p_\theta(\pi_i | s_i)$ 
9:      $\mathcal{L}_v \leftarrow \frac{1}{B} \sum_{i=1}^B \|b_i - L(\pi_i)\|_2^2$ 
10:     $\theta \leftarrow \text{ADAM}(\theta, g_\theta)$ 
11:     $\theta_v \leftarrow \text{ADAM}(\theta_v, \nabla_{\theta_v} \mathcal{L}_v)$ 
12:  end for
13:  return  $\theta$ 
14: end procedure
```

Pseudo code of actor-critic training

● Search strategies

- RL pretraining - Greedy
 - After training, select the tours from $p_\theta(\cdot | s)$ and greedy method
- RL pretraining - Sampling
 - After training, select the multiple candidate tours(1.28M) from stochastic policy $p_\theta(\cdot | s)$ and select the shortest one
 - Controlling the diversity of the sampled tour with a temperature hyperparameter T
 - $A(ref, q, T; W_{ref}, W_q, v) \triangleq softmax(u/T)$
- Active Search
 - Refining the parameters of stochastic policy $p_\theta(\cdot | s)$ during inference to minimize $J(\theta)$ on a single test input \mathbf{s}
 - Only policy updated, not critic.
- RL pretraining - Active Search
 - After training, using active search

Algorithm 2 Active Search

```

1: procedure ACTIVESEARCH(input  $s, \theta$ , number of candidates  $K, B, \alpha$ )
2:    $\pi \leftarrow \text{RANDOMSOLUTION}()$ 
3:    $L_\pi \leftarrow L(\pi | s)$ 
4:    $n \leftarrow \lceil \frac{K}{B} \rceil$ 
5:   for  $t = 1 \dots n$  do
6:      $\pi_i \sim \text{SAMPLESOLUTION}(p_\theta(\cdot | s))$  for  $i \in \{1, \dots, B\}$ 
7:      $j \leftarrow \text{ARGMIN}(L(\pi_1 | s) \dots L(\pi_B | s))$ 
8:      $L_j \leftarrow L(\pi_j | s)$ 
9:     if  $L_j < L_\pi$  then
10:       $\pi \leftarrow \pi_j$ 
11:       $L_\pi \leftarrow L_j$ 
12:     end if
13:      $g_\theta \leftarrow \frac{1}{B} \sum_{i=1}^B (L(\pi_i | s) - b) \nabla_\theta \log p_\theta(\pi_i | s)$ 
14:      $\theta \leftarrow \text{ADAM}(\theta, g_\theta)$ 
15:      $b \leftarrow \alpha \times b + (1 - \alpha) \times (\frac{1}{B} \sum_{i=1}^B b_i)$ 
16:   end for
17:   return  $\pi$ 
18: end procedure

```

Pseudo code of active search

Configuration	Learn on training data	Sampling on test set	Refining on test set
RL pretraining-Greedy	Yes	No	No
Active Search (AS)	No	Yes	Yes
RL pretraining-Sampling	Yes	Yes	No
RL pretraining-Active Search	Yes	Yes	Yes

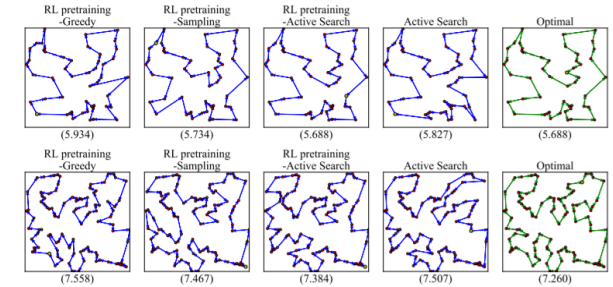
Different learning configurations

Experiment Results

Experiment results (I)

● Comparison with different strategies for TSP family

- 3 개의 TSP에 대해서, 논문에서 제시한 기법과, 기존에 사용되던 기법들, 최적 해 사이의 성능을 비교하는 실험을 수행(TSP20, TSP50, TSP100)
 - Supervised learning(Pointer network)
 - RL pretraining - greedy
 - RL pretraining - sampling
 - RL pretraining - active search
 - Active search
 - Christofides: polynomial time complexity
 - OR Tools' local search: made by Google
 - Optimal solution
- 각 기법들을 비교할 성능 지표로는, Average tour length와 running time을 선택
 - RL+sampling, RL+AS의 경우, optimal에 근접한 결과를 얻어냈음을 확인
 - RL+greedy의 경우, 다른 기법들과 비교해 보았을 때 , 시간 효율적임을 확인



Sample results. TSP50, TSP100

Task	Supervised Learning	RL pretraining				AS	Christo-fides	OR Tools' local search	Optimal
		greedy	greedy@16	sampling	AS				
TSP20	3.88 ^(†)	3.89	—	3.82	3.82	3.96	4.30	3.85	3.82
TSP50	6.09 ^(†)	5.95	5.80	5.70	5.70	5.87	6.62	5.80	5.68
TSP100	10.81	8.30	7.97	7.88	7.83	8.19	9.18	7.99	7.77

Average tour length of different methods

Task	RL pretraining		OR-Tools' local search	Optimal	
	greedy	greedy@16		Concorde	LK-H
TSP50	0.003s	0.04s	0.02s	0.05s	0.14s
TSP100	0.01s	0.15s	0.10s	0.22s	0.88s

Running time in seconds

Experiment results (II)

● Comparison with different strategies for KnapSack problems

- 3 개의 KnapSack problem에 대해서, 논문에서 제시한 기법 일부와, 기존에 사용되던 기법들, 최적 해 사이의 성능을 비교하는 실험을 수행(KNAP50, KNAP100, KNAP200)
 - RL pretraining - greedy
 - Active Search
 - Random search
 - Greedy
 - Optimal solution
- 각 기법들을 비교할 성능 지표로는, total value z 로 설정
 - RL+greedy의 경우, optimal solution의 1% 이하의 오차율을 보였음을 확인
 - Active Search의 경우, 최적해를 찾았음을 확인

Task	RL pretraining greedy	Active Search	Random Search	Greedy	Optimal
KNAP50	19.86	20.07	17.91	19.24	20.07
KNAP100	40.27	40.50	33.23	38.53	40.50
KNAP200	57.10	57.45	35.95	55.42	57.45

Total value of different methods

Thank you!

Q&A

Implementation of NCO

● Environment: OR-gym

■ Knapsack-v0

- 무게 제한 아래에서, 최고의 효용을 얻을 수 있는 물품들을 가방에 담는 문제
- $z = \sum_{i=1}^n v_i x_i$, s.t. $\sum_{i=1}^n w_i x_i \leq W$, $x_i \in \mathbb{Z}^{0+}$

■ MDP formulation

- State:
 - [Item weights, item values, current weight]
- Action:
 - Number of item
- Reward:
 - Item value : value of item
 - Over packed penalty : -100

■ TSP-v0

- 최소 비용으로 모든 지역을 순회하는 방법에 관해 묻는 문제
- $z = ||x_n - x_1||_2 + \sum_{i=1}^{n-1} ||x_i - x_{i+1}||_2$

■ MDP formulation

- State:
 - [current node, connection(adjacency matrix)]
- Action:
 - Number of node
- Reward:
 - Cost of move : -distance
 - Invalid action penalty : -100
 - Success : 1000

- **LunarLanderContinuous-v2**

- Comparison with Vanilla TD3(**TD3_None**), TD3_gSDE(8, 16, 32, **64, Episode**)
 - Figure 1 shows that TD3_gSDE is better stability than Vanilla TD3

실험 중

Learning environment