

Diffuser

윤승제

Planning with Diffusion for Flexible Behavior Synthesis

Michael Janner^{*1} Yilun Du^{*2} Joshua B. Tenenbaum² Sergey Levine¹

Abstract

Model-based reinforcement learning methods often use learning only for the purpose of estimating an approximate dynamics model, offloading the rest of the decision-making work to classical trajectory optimizers. While conceptually simple, this combination has a number of empirical shortcomings, suggesting that learned models may not be well-suited to standard trajectory optimization. In this paper, we consider what it would look like to fold as




Figure 1. Diffuser is a diffusion probabilistic model that plans by iteratively refining trajectories.

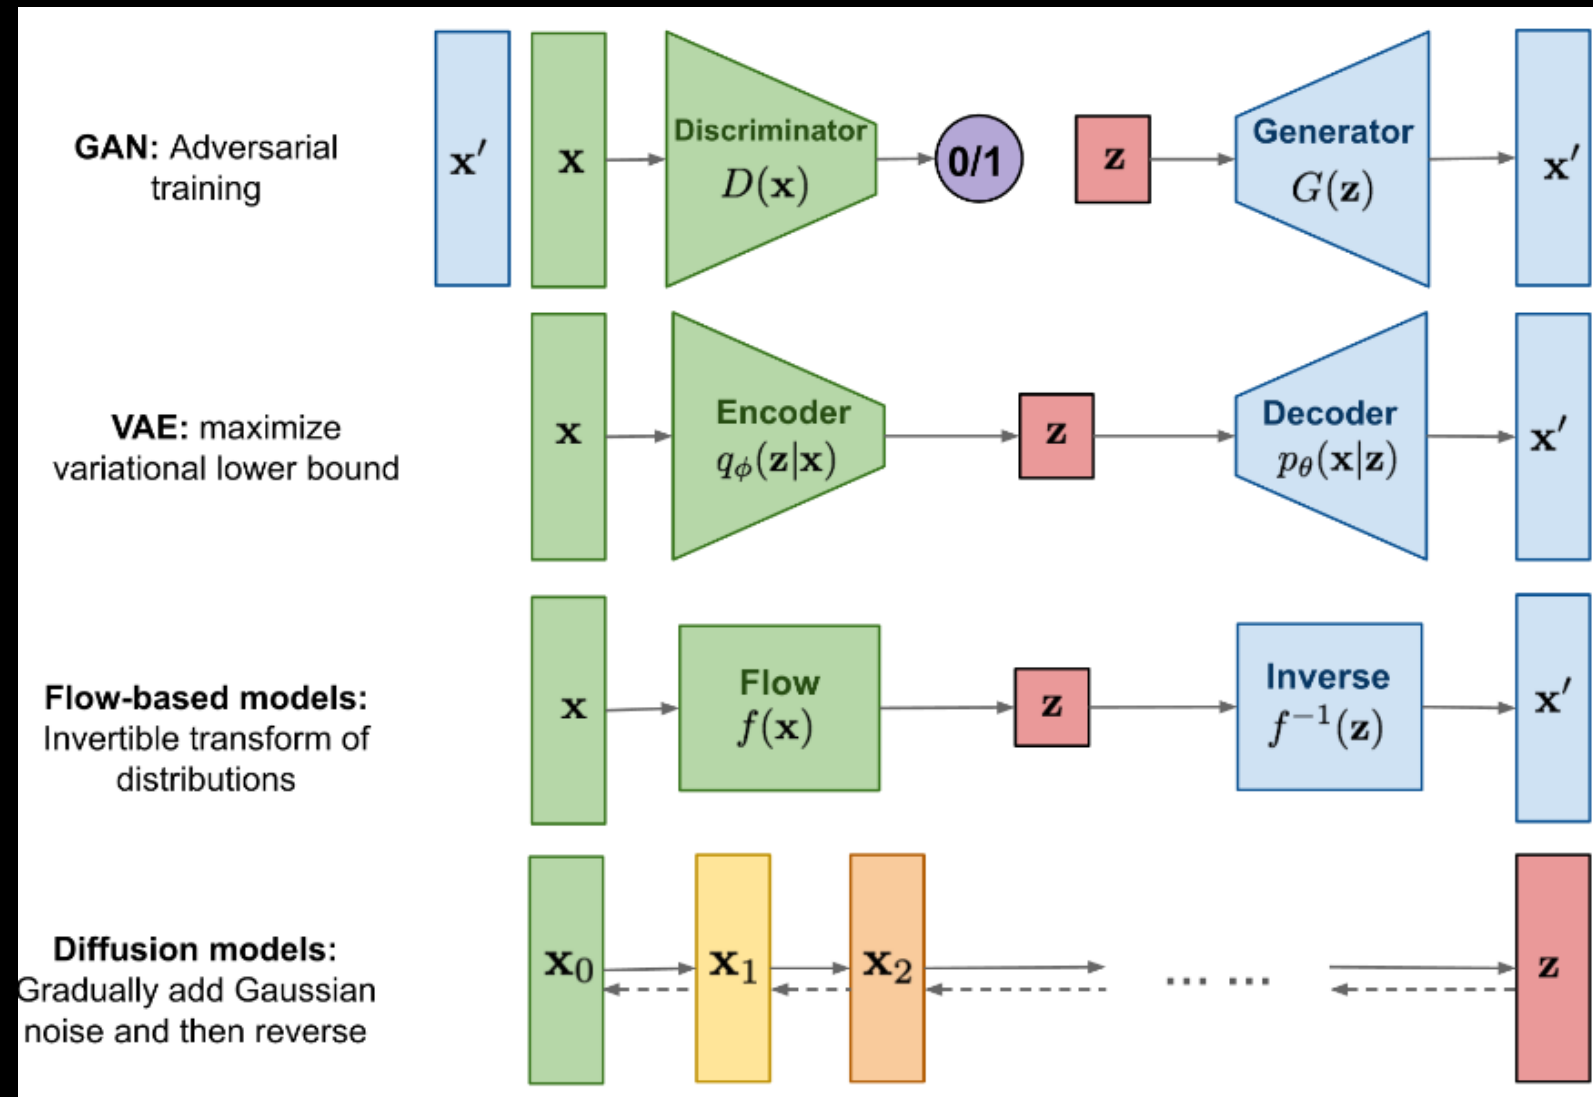
Contents

- Background
 - Denoising diffusion probabilistic model
 - Offline Reinforcement Learning
 - Model based Reinforcement learning
- Diffuser
 - Structure
 - Guided Sampling for planning
 - Properties
- Results
 - Multi task planning
 - Test time flexibility
 - conclusion

Background

Overview of DDPM

- Generative model



<https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>

Denoising Diffusion Probabilistic Models

Denoising Diffusion Probabilistic Models

Jonathan Ho

UC Berkeley

`jonathanho@berkeley.edu`

Ajay Jain

UC Berkeley

`ajayj@berkeley.edu`

Pieter Abbeel

UC Berkeley

`pabbeel@cs.berkeley.edu`

Overview of DDPM

- Denoising diffusion probabilistic model
 - Forward q : noise 추가
 - Reverse p : noise 에서 원본 이미지로 denoising




Figure 2: The directed graphical model considered in this work.

Denoising Diffusion Probabilistic Models 2020

Overview of DDPM


- Denoising diffusion probabilistic model
 - 본래 noise 공간에서 large scale 먼저 거시적인 부분부터 생성
 - 거시적인 요소들이 생성되고 난 후 점차 미세적인 부분도 구체화되면서 생성



Figure 6: Unconditional CIFAR10 progressive generation (\hat{x}_0 over time, from left to right). Extended samples and sample quality metrics over time in the appendix (Figs. 10 and 14).

Overview of DDPM

- Interpolation of DDPM
 - 픽셀 상에서의 interpolation 이 아닌 diffusion 된 space 상에서 interpolation



Training process of DDPM

- How to train the forward/reverse process
 - 최종 프로세스

Algorithm 1 Training

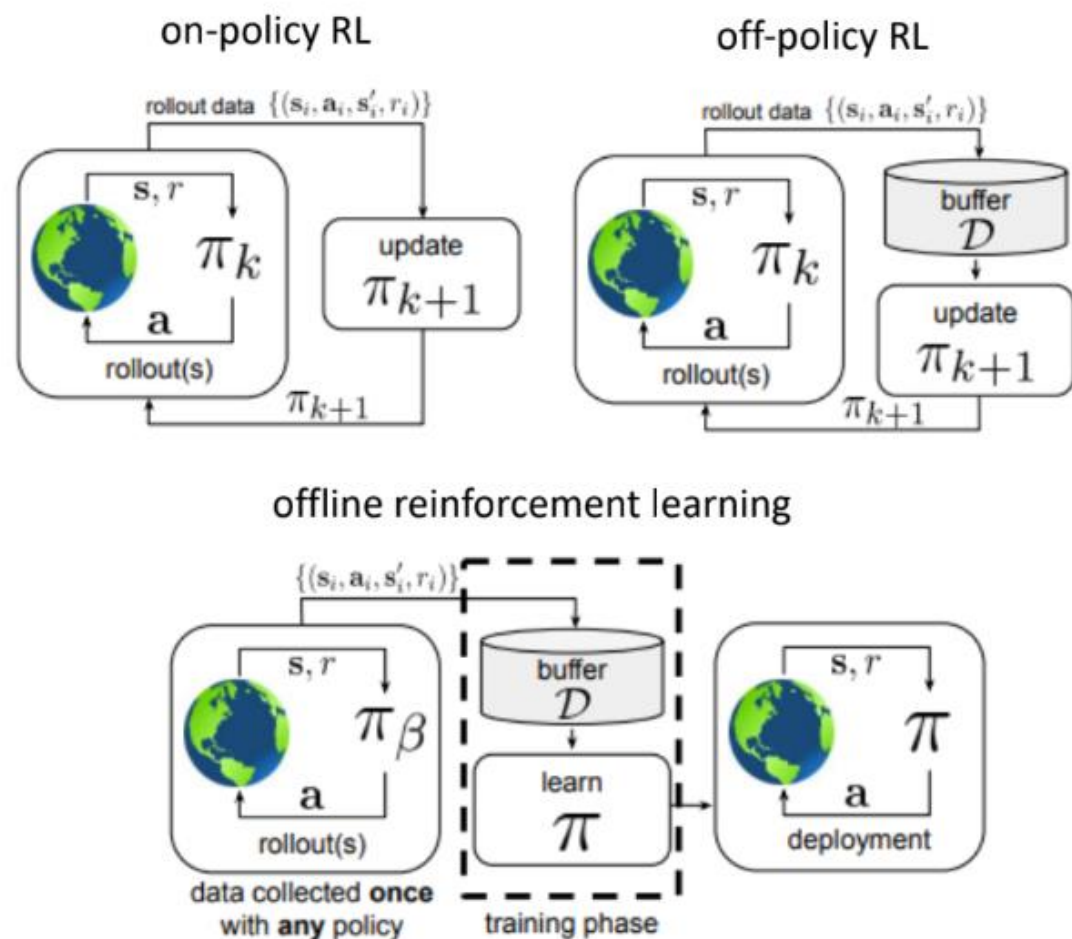
```
1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$ 
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take gradient descent step on
        $\nabla_{\theta} \|\epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2$ 
6: until converged
```

Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```

Offline RL

- Training process of offline RL
 - Sim, env와의 interaction 이 없음
 - Buffer 내에 없는 Counterfactual 한 scene에 대응하기 위해 uncertainty estimation & Q regularization 에 집중



Formally:

$$\mathcal{D} = \{(s_i, a_i, s'_i, r_i)\}$$

$$s \sim d^{\pi_\beta}(s)$$

$$a \sim \pi_\beta(a|s)$$

$$s' \sim p(s'|s, a)$$

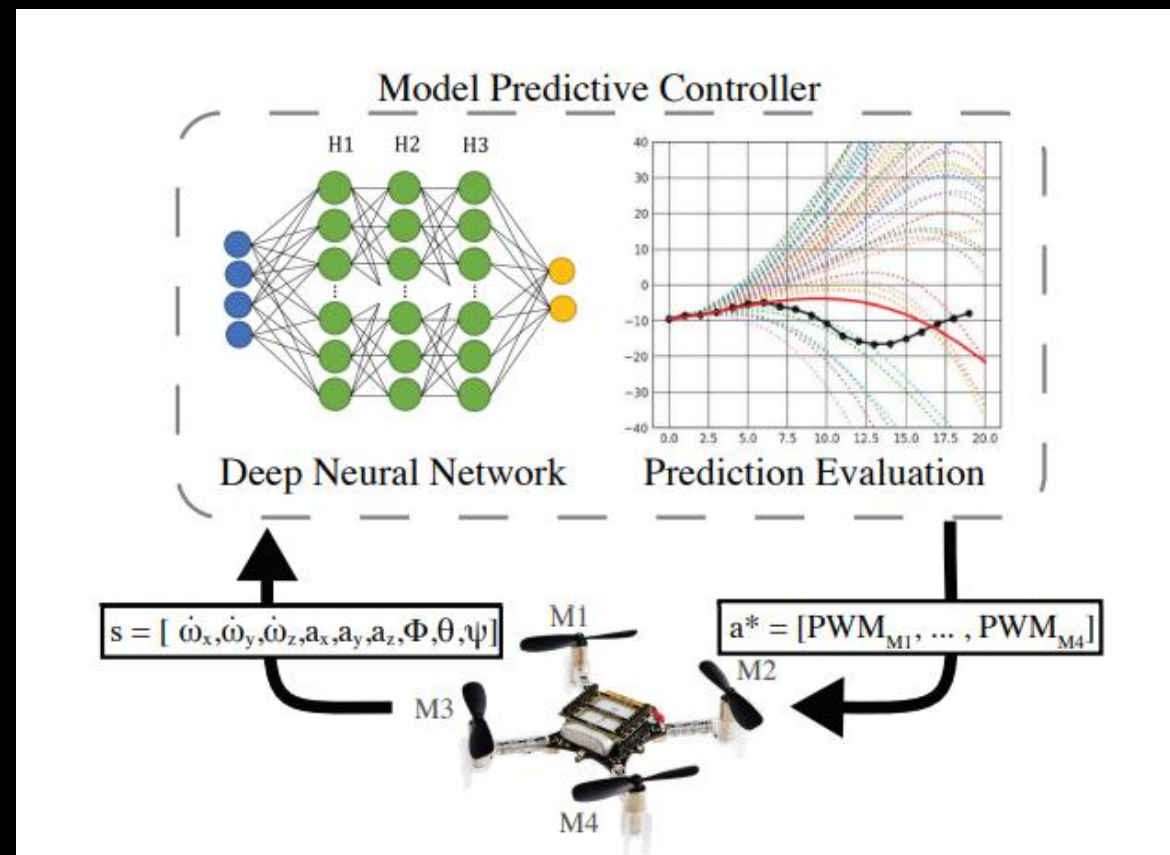
$$r \leftarrow r(s, a)$$

← generally **not** known

$$\text{RL objective: } \max_{\pi} \sum_{t=0}^T E_{s_t \sim d^{\pi}(s), a_t \sim \pi(a|s)} [\gamma^t r(s_t, a_t)]$$

Model based RL

- Examples of model based RL
 - Transition model을 supervised learning 하는 것이 일반적
 - Scalable for several task and sample efficient
 - 기존 one step prediction + markovian이라 Long term prediction 을 신뢰하긴 어려움
 - Planning 과정에서 생성된 trajectory가 실제 환경의 trajectory space에 adversarial 할수도



Rambert et al, "Low Level Control of a Quadrotor with Deep Model-Based Reinforcement Learning", IROS19

Diffuser : DDPM based trajectory synthesis

Diffuser

- Intuition

- Janner 왈 : Transition model sl & predictive planning을 generative model로 대체?

- Offload as much of MBRL into contemporary generative modeling as possible

replace **prediction and planning** with big generative model

Algorithm 1 Model-based RL (idealized)

Inputs: Dataset of transitions $\mathcal{D} = \{(s_t, \mathbf{a}_t, s_{t+1}), \dots\}$, reward function $r(\cdot, \cdot)$, current state s_0

1: Train a **predictive model**


$$\underset{f}{\text{minimize}} \mathbb{E}_{\mathbf{s}_t, \mathbf{a}_t, \mathbf{s}_{t+1} \sim \mathcal{D}} [\|\mathbf{s}_{t+1} - f(\mathbf{s}_t, \mathbf{a}_t)\|]$$

2: Use model to evaluate potential plans $\mathbf{a}_{0:T}$, selecting the best one:

$$\underset{\mathbf{a}_{0:T}}{\text{maximize}} r(\mathbf{s}_0, \mathbf{a}_0) + r(\mathbf{s}_1, \mathbf{a}_1) + r(\mathbf{s}_2, \mathbf{a}_2) + \dots$$

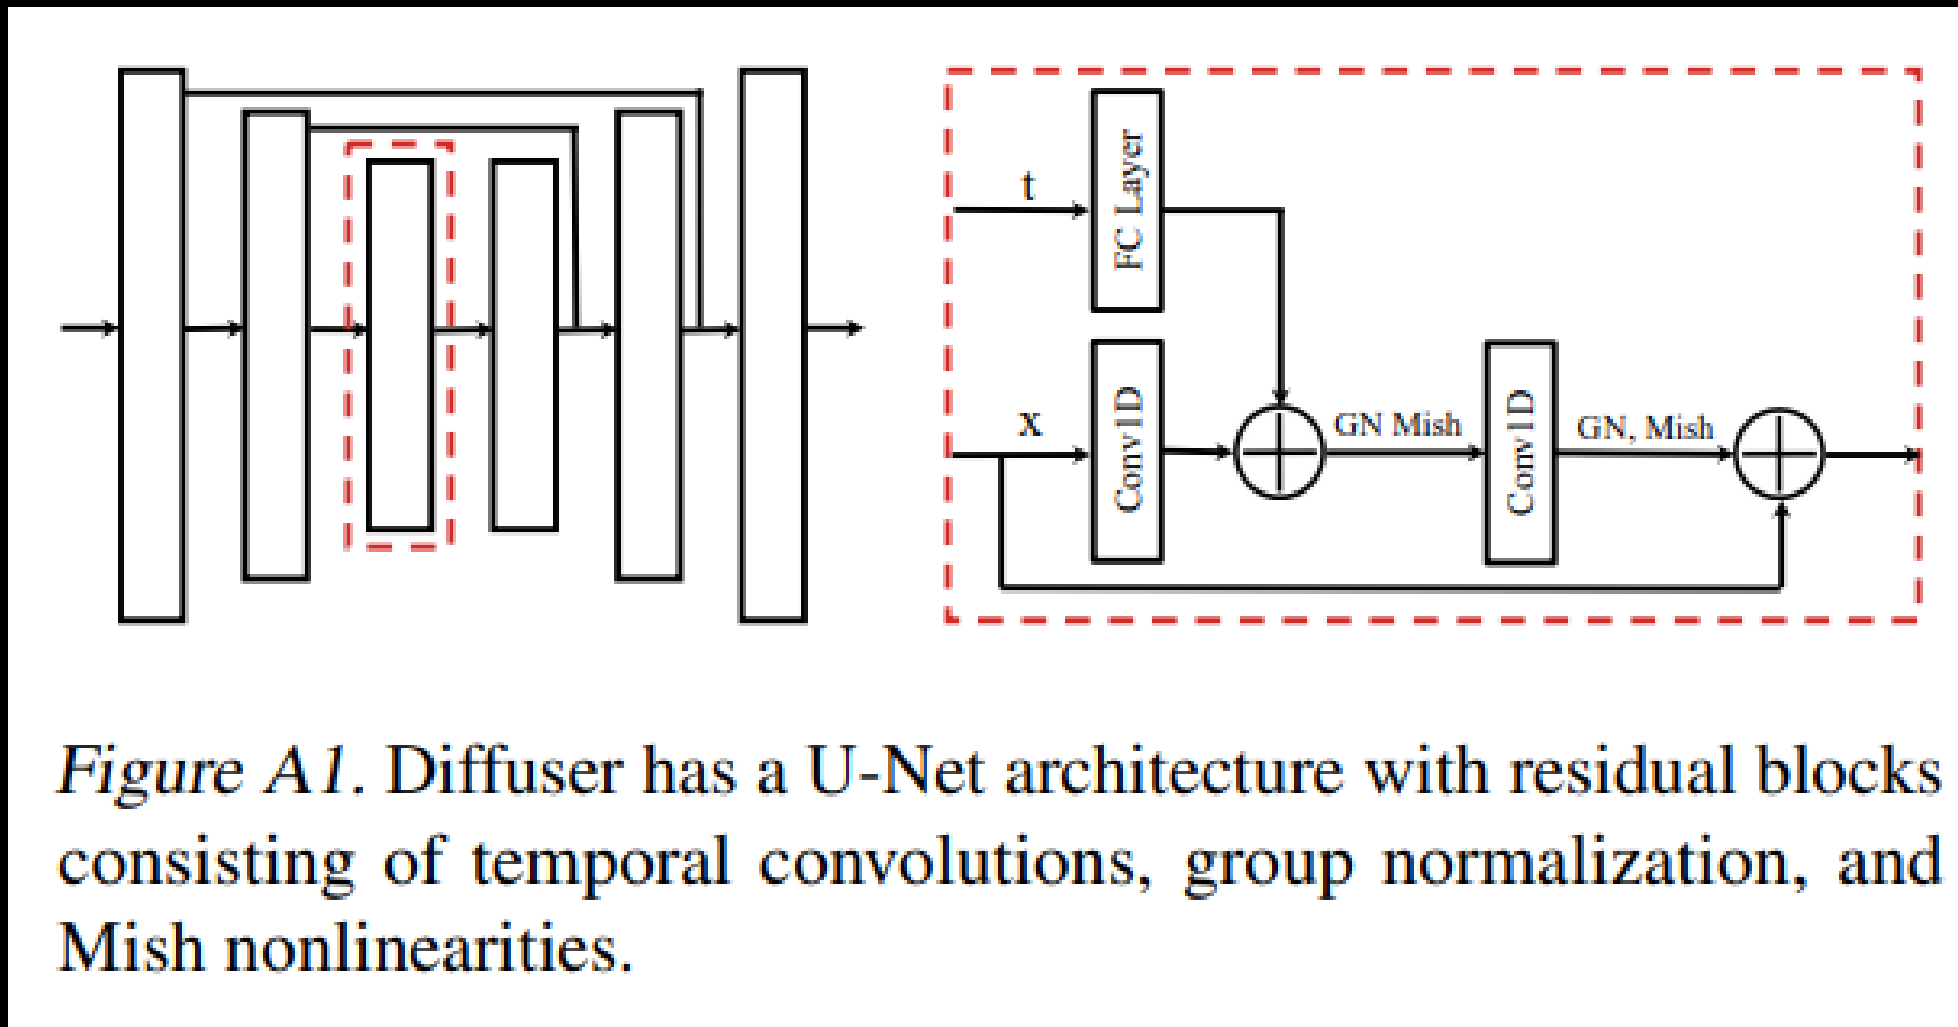
Diffuser

- Structure of diffuser
 - Single channel image array 처럼 s , a trajectory noise array 로 초기화.
 - Diffuser를 통해서 reward maximizing 한 방향으로 denoising.
 - Not used rnn but local receptive field - > rnn 특유의 memoryless 문제에 자유롭다.



Diffuser

- Structure of diffuser
 - Unet 구조
 - Group Norm
 - Mish activation



Diffuser

- Planning with diffusion
 - 아래 식처럼 trajectory에 대한 분포를 정의

$$\tilde{p}_{\theta}(\tau) \propto p_{\theta}(\tau)h(\tau).$$


- $p_{\theta}(\tau)$: physically realistic
 - $h(\tau)$: guidance (higher rewards)
- Guided sampling
 - $h(\tau) = p(O_{1:t}|\tau)$, O_i 는 binary, optimal이면 1 아니면 0
 - $p(O_t = 1) = \exp(r(s_t, a_t))$
- Goal conditioned RL as inpainting

$$h(\tau) = \delta_{\mathbf{c}_t}(\mathbf{s}_0, \mathbf{a}_0, \dots, \mathbf{s}_T, \mathbf{a}_T) = \begin{cases} +\infty & \text{if } \mathbf{c}_t = \mathbf{s}_t \\ 0 & \text{otherwise} \end{cases}$$

Diffuser

- Guided sampling

- Guide operator \mathcal{J} 는 0부터 T 까지 trajectory에서 얻은 return 총합
- \mathcal{J} 를 최대화 하는 방향으로 gradient를 구하고, denoising 하는 reverse step p 에서 mean을 guide



$$p_{\theta}(\tau^{i-1} | \tau^i, \mathcal{O}_{1:T}) \approx \mathcal{N}(\tau^{i-1}; \mu + \Sigma g, \Sigma)$$

$$\begin{aligned} g &= \nabla_{\tau} \log p(\mathcal{O}_{1:T} | \tau) |_{\tau=\mu} \\ &= \sum_{t=0}^T \nabla_{\mathbf{s}_t, \mathbf{a}_t} r(\mathbf{s}_t, \mathbf{a}_t) |_{(\mathbf{s}_t, \mathbf{a}_t)=\mu_t} = \nabla \mathcal{J}(\mu). \end{aligned}$$

Algorithm 1 Guided Diffusion Planning

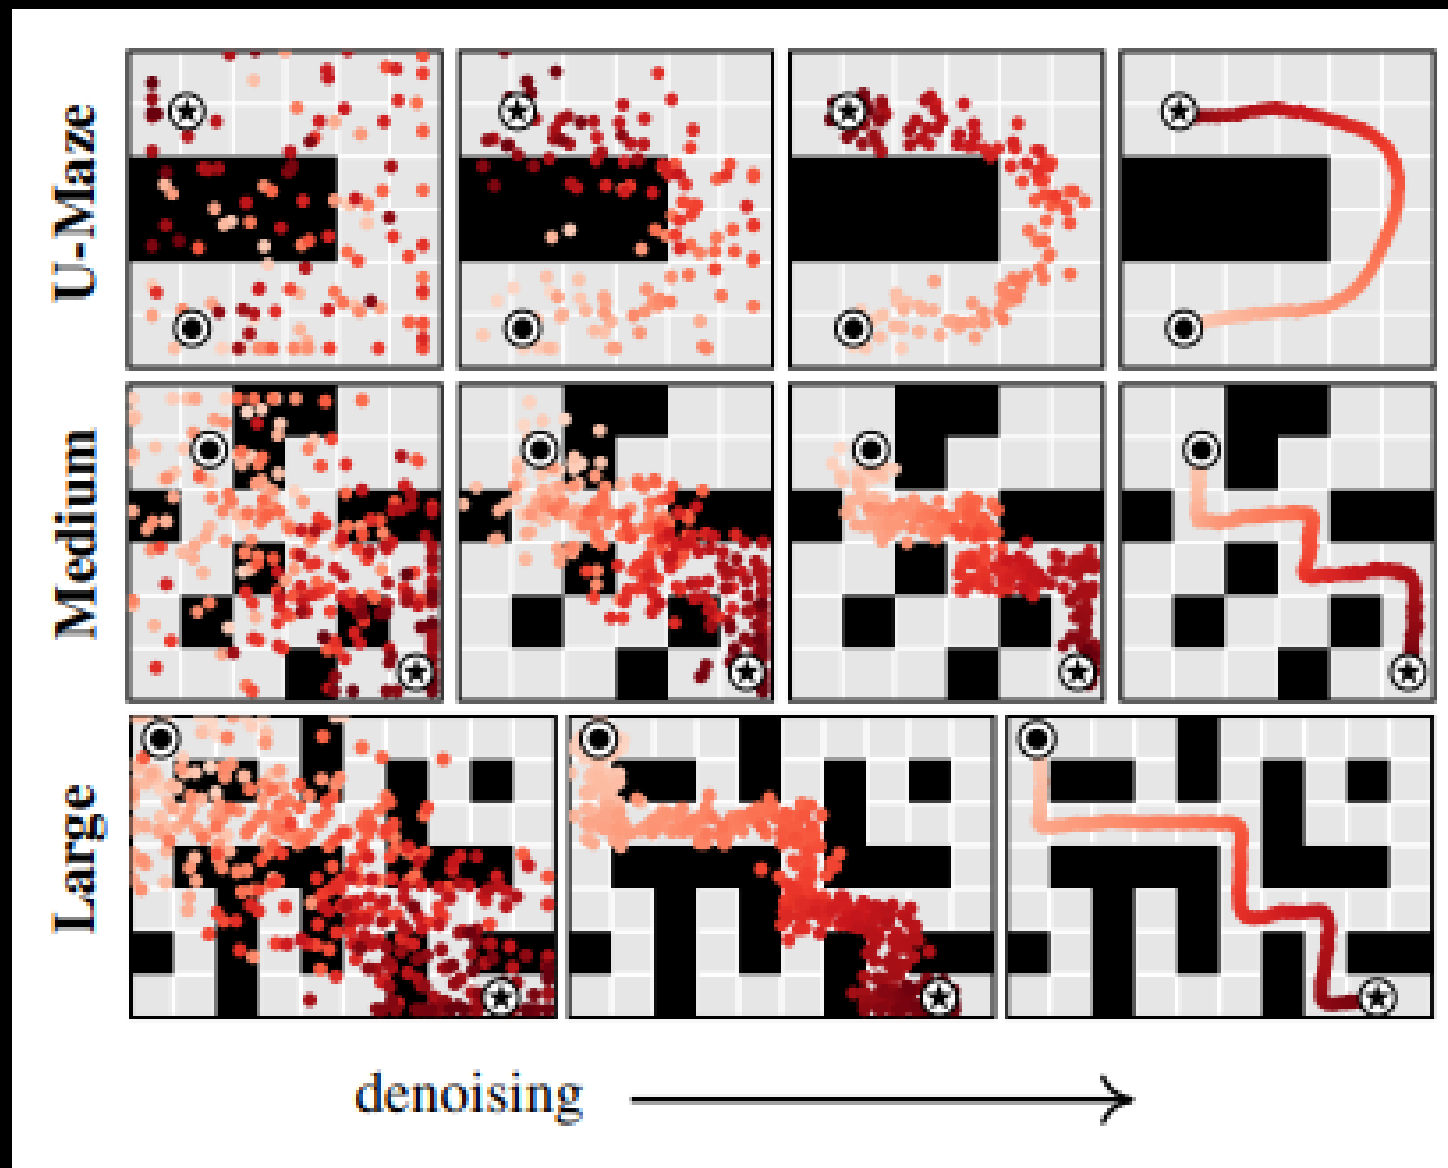
```

1: Require Diffuser  $\mu_{\theta}$ , guide  $\mathcal{J}$ , scale  $\alpha$ , covariances  $\Sigma^i$ 
2: while not done do
3:   Observe state  $\mathbf{s}$ ; initialize plan  $\tau^N \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
4:   for  $i = N, \dots, 1$  do
5:     // parameters of reverse transition
6:      $\mu \leftarrow \mu_{\theta}(\tau^i)$ 
7:     // guide using gradients of return
8:      $\tau^{i-1} \sim \mathcal{N}(\mu + \alpha \Sigma \nabla \mathcal{J}(\mu), \Sigma^i)$ 
9:     // constrain first state of plan
10:     $\tau_{s_0}^{i-1} \leftarrow \mathbf{s}$ 
11:   Execute first action of plan  $\tau_{\mathbf{a}_0}^0$ 

```


Properties of diffuser

- Learned long-horizon planning
 - 기존 model based RL과 달리 single step으로 dynamics를 예측하지 않고, horizon T 만큼의 noise array를 denoising 해서 기존보다 long-term prediction 을 효율적으로 처리
 - Starting points와 goal을 다르게 세팅해도 trajectory 생성 가능
 - Sparse reward 환경에서 강점을 가진다.




Properties of diffuser

- Temporal compositionality
 - Autoregressive 한 요소 (rnn) 없고, non-markovian 이지만, denoising 과정에서 계속 local receptive field로 temporal 한 부분을 refine 해주는 방식이라, markovian 처럼 거동 가능
 - 데이터 내에서 존재하는 temporal, transition 특성들을 학습해서 그 특성 안에 feasible한 trajectory 로 생성




data



plan

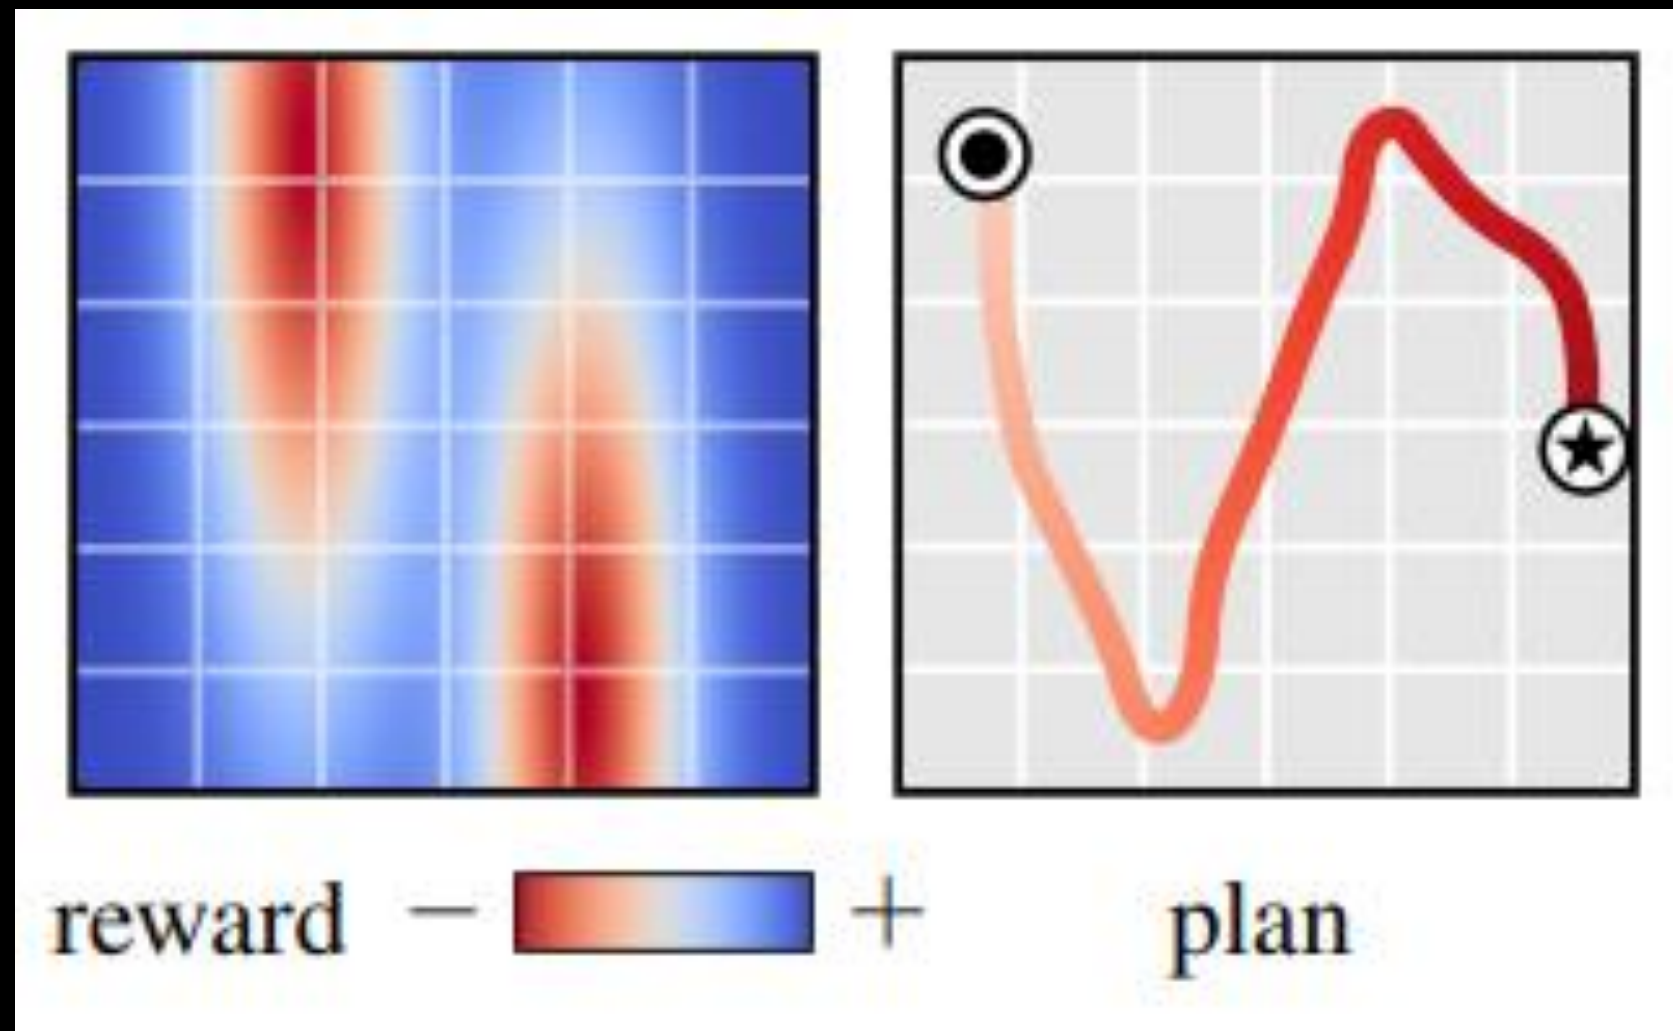
Properties of diffuser

- Variable length planning
 - Latent noise array의 길이에 따라서 생성되는 trajectory의 길이가 결정됨



Properties of diffuser

- Task compositionality
 - 새로운 reward function 으로 부터 학습 과정에 들어가지 않았던 새 task에 대해 composition 이 가능



Results of Diffuser

Results of diffuser


- Multi task planning
 - Baseline인 BCQ, CQL, IQL 보다 좋은 성능
 - BCQ, CQL, IQL 모두 value network에 goal condition을 input으로 줘야하는 형식이라 multi task 시 학습을 다시 해야 하나, Diffuser는 그럴 필요 없음.
 - Denoing 과정에서 s_0, s_T 을 원하는 값으로 계속 fix 해주면 multi task 에도 바로 사용 가능.

Environment		BCQ	CQL	IQL	Diffuser
Maze2D	U-Maze	12.8	5.7	47.4	113.9 ± 3.1
Maze2D	Medium	8.3	5.0	34.9	121.5 ± 2.7
Maze2D	Large	6.2	12.5	58.6	123.0 ± 6.4
Single-task Average		9.1	7.7	47.0	119.5
Multi2D	U-Maze	-	-	24.8	128.9 ± 1.8
Multi2D	Medium	-	-	12.1	127.2 ± 3.4
Multi2D	Large	-	-	13.9	132.1 ± 5.8
Multi-task Average		-	-	16.9	129.4

Results of diffuser

- Test time flexibility

- 환경은 block stacking
- 학습 과정에 없었던 새로운 task에서도 잘 되느냐
- 100이 perfect
- 학습에 없었던 새로운 task에 전혀 동작을 못하는 BCQ, CQL에 비해 diffuser는 40-60 사이의 성능을 보임 (중간에 constraint가 틀렸다면, task 가 완료가 안 되서 감점된 듯)



Environment	BCQ	CQL	Diffuser
Unconditional Stacking	0.0	24.4	58.7 ± 2.5
Conditional Stacking	0.0	0.0	45.6 ± 3.1
Rearrangement	0.0	0.0	58.9 ± 3.4
Average	0.0	8.1	54.4

Reference

- Reference

- <https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>
- <https://arxiv.org/abs/2006.11239>
- <https://diffusion-planning.github.io/>
- <https://arxiv.org/abs/1901.03737>

END