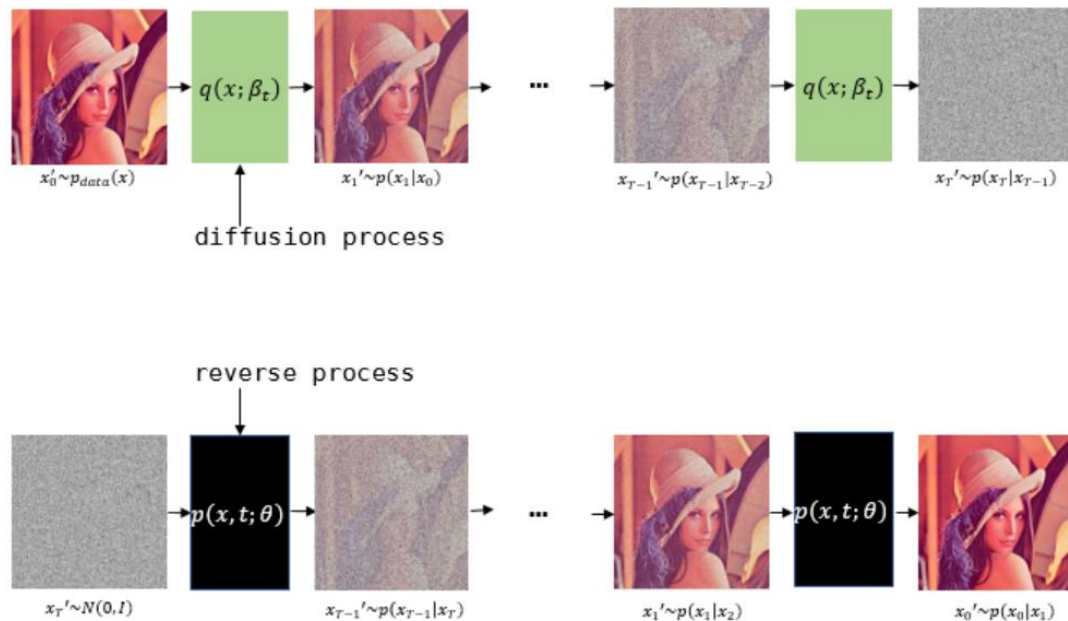



Diffusion Policies as an Expressive Policy Class for Offline Reinforcement Learning

Related works : Diffusion model(DDPM)



- Diffusion process에서는 원래 이미지에 noise를 조금씩, 반복적으로 추가하여 원래 이미지와 거의 independent한 noise인 latent variable을 생성
- 각 step은 Markov decision process라 가정
- Reverse process는 이 diffusion process의 역과정을 neural network를 통해 학습

Diffusion model : Diffusion process



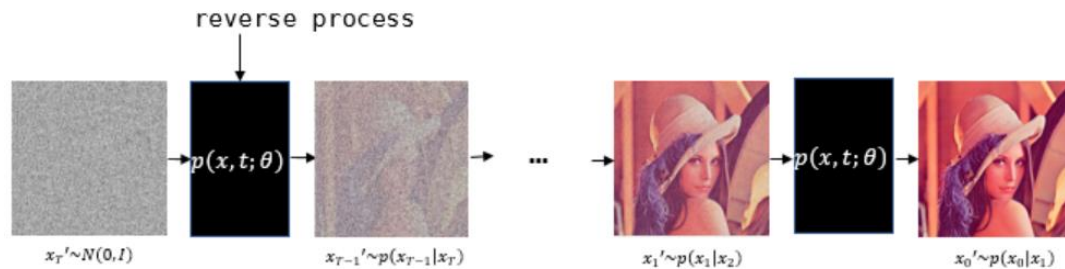
$$q(x_{1:T}|x_0) = \prod_{t=1}^T q(x_t|x_{t-1}), \quad q(x_t|x_{t-1}) = N(\sqrt{1 - B_t}x_{t-1}, B_t I)$$

- q 함수는 diffusion process이며 미세한 gaussian noise를 추가하는 과정
- Diffusion process는 trainable parameter가 없음(VDM에서 trainable로 변경)
- B_t 는 noise scheduler

$$\begin{aligned} q(x_t|x_0) &= N(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I) \\ &= \sqrt{\bar{\alpha}_t} x_0 + \varepsilon \sqrt{1 - \bar{\alpha}_t}, \varepsilon \sim N(0, I) \\ &\text{with } \alpha_t = 1 - B_t \text{ and } \bar{\alpha}_t = \prod_{s=0}^t \alpha_s \end{aligned}$$

- 위의 수식으로 t time에 있는 latent variable을 바로 계산할 수 있음

Diffusion model : Reverse process



$$q(x_t | x_{t-1}) \rightarrow q(x_{t-1} | x_t)$$

$$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T p_\theta(x_{t-1} | x_t)$$
$$p_\theta(x_{t-1} | x_t) = N(\mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

- mean에 해당하는 μ_θ 와 covariance Σ_θ 학습이 목적
- NF에서는 change of variable theorem 으로 계산했지만 diffusion은 neural network를 사용
- 이때 diffusion process와 마찬가지로 reverse process 또한 Markov chain으로 가정함

Diffusion model (DDPM)

- DDPM의 핵심은 neural network로 표현되는 p 함수가 q 를 보고 noise를 걷어내는 과정을 학습하는 것
- loss는 VAE의 loss와 유사하게 negative log-likelihood로 전개됨

$$\begin{aligned} LOSS_{Diffusion} &= D_{KL}(q(z | x_0) || P_{\theta}(x_0 | z)) - E_{z \sim q(z|x)}[\log P_{\theta}(z)] \\ &= D_{KL}(q(z | x_0) || P_{\theta}(z)) + \sum_{t=2} D_{KL}(q(x_{t-1} | x_t, x_0) || P_{\theta}(x_{t-1} | x_t)) - E_q[\log P_{\theta}(x_0 | x_1)] \end{aligned}$$

Regularizer on Encoder *Denoising Process* *Reconstruction on Decoder*



DDPM에서는 Loss가 굉장히 간단한 식으로 정의됨

$$LOSS_{DDPM} = \mathbb{E}_{x_0, \epsilon} \left[\left| \epsilon - \epsilon_{\theta} \left(\sqrt{\tilde{\alpha}_t} + \sqrt{1 - \tilde{\alpha}_t} \epsilon, t \right) \right|^2 \right]$$

Related works : Offline RL

1. Constraining the learned value function to assign low values to OOD actions
2. Introducing model-based methods, which learn a model of the environment dynamics and perform pessimistic planning in the learned MDP
3. Treating offline RL as a problem of sequence prediction with return guidance(Offline Reinforcement Learning as One Big Sequence Modeling Problem,2021)
4. Regularizing how far the policy can deviate from the behavior policy

A method to perform policy regularization using diffusion models

Preliminaries

- MDP : $M = \{S, A, P, R, \gamma, d_0\}$
- In the offline setting, a static dataset $D \triangleq \{(s, a, r, s')\}$
- Diffusion process : $q(x_{1:T}|x_0) := \prod_{t=1}^T q(x_t|x_{t-1})$
- Reverse diffusion chain, $p_\theta(x_{0:T}) := N(x_T; 0, I) \prod_{t=1}^T p_\theta(x_{t-1}|x_t)$

Diffusion policy

- RL policy via the reverse process of a conditional diffusion model

$$\pi_{\theta}(a|s) = p_{\theta}(a^{0:N}|s) = N(a^N; 0, I) \prod_{t=1}^T p_{\theta}(a^{i-1}|a^i, s)$$

- We first sample $a^N \sim N(0, I)$ and then from the reverse diffusion chain as

$$p_{\theta}(a^{i-1}|a^i) = \frac{a^i}{\sqrt{\alpha_i}} - \frac{\beta_i}{\sqrt{1-\alpha_i}} \epsilon_{\theta}(a^i, s^i, i) + \sqrt{\beta_i} \epsilon$$
$$\beta_i = 1 - \alpha_i = 1 - e^{-\beta_{\min}(\frac{1}{N}) - 0.5(\beta_{\max} - \beta_{\min})\frac{2i-1}{N^2}},$$

- Objective function is proposed by DDPM,

$$\mathcal{L}_d(\theta) = \mathbb{E}_{i \sim \mathcal{U}, \epsilon \sim \mathcal{N}(\mathbf{0}, I), (s, a) \sim \mathcal{D}} [\|\epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_i} a + \sqrt{1 - \bar{\alpha}_i} \epsilon, s, i)\|^2]$$

Q-learning

- The policy-regularization loss $L_d(\theta)$ is a behavior-cloning term
- To improve the policy, we inject Q-value function guidance into the reverse diffusion chain in the training stage in order to learn to preferentially sample actions with high values

$$\pi = \arg \min_{\pi_\theta} \mathcal{L}(\theta) = \mathcal{L}_d(\theta) + \mathcal{L}_q(\theta) = \mathcal{L}_d(\theta) - \alpha \cdot \mathbb{E}_{\mathbf{s} \sim \mathcal{D}, \mathbf{a}^0 \sim \pi_\theta} [Q_\phi(\mathbf{s}, \mathbf{a}^0)]$$

- We build two Q-networks, and target networks, optimize formula as

$$\mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t, \mathbf{s}_{t+1}) \sim \mathcal{D}, \mathbf{a}_{t+1}^0 \sim \pi_{\theta'}} \left[\left\| (r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \min_{i=1,2} Q_{\phi'_i}(\mathbf{s}_{t+1}, \mathbf{a}_{t+1}^0)) - Q_{\phi_i}(\mathbf{s}_t, \mathbf{a}_t) \right\|^2 \right].$$

Algorithm

Algorithm 1 Diffusion Q-learning

Initialize policy network π_θ , critic networks Q_{ϕ_1} and Q_{ϕ_2} , and target networks $\pi_{\theta'}$, $Q_{\phi'_1}$ and $Q_{\phi'_2}$

for each iteration **do**

Sample transition mini-batch $\mathcal{B} = \{(s_t, a_t, r_t, s_{t+1})\} \sim \mathcal{D}$.

Q-value function learning

Sample $a_{t+1}^0 \sim \pi_{\theta'}(a_{t+1} | s_{t+1})$ by Equation (1). $p_\theta(a^{i-1} | a^i) = \frac{a^i}{\sqrt{\alpha_i}} - \frac{\beta_i}{\sqrt{1-\alpha_i}} \varepsilon_\theta(a^i, s^i, i) + \sqrt{\beta_i} \varepsilon$

Update Q_{ϕ_1} and Q_{ϕ_2} by Equation (4). (max Q backup by Kumar et al. (2020) could be added)

$$\mathbb{E}_{(s_t, a_t, s_{t+1}) \sim \mathcal{D}, a_{t+1}^0 \sim \pi_{\theta'}} \left[\left\| (r(s_t, a_t) + \gamma \min_{i=1,2} Q_{\phi'_i}(s_{t+1}, a_{t+1}^0)) - Q_{\phi_i}(s_t, a_t) \right\|^2 \right].$$

Policy learning

Sample $a_t^0 \sim \pi_\theta(a_t | s_t)$ by Equation (1).

Update policy by minimizing Equation (3). $\pi = \underset{\pi_\theta}{\operatorname{argmin}} L(\theta) = L_d(\theta) + L_q(\theta)$

Update target networks

$\theta' = \rho\theta' + (1 - \rho)\theta$, $\phi'_i = \rho\phi'_i + (1 - \rho)\phi_i$ for $i = \{1, 2\}$.

end for

Experiments

Gym Tasks	BC	AWAC	Diffuser	MoRel	Onestep RL	TD3+BC	DT	CQL	IQL	Diffusion-QL
halfcheetah-medium-v2	42.6	43.5	44.2	42.1	48.4	48.3	42.6	44.0	47.4	51.1 \pm 0.5
hopper-medium-v2	52.9	57.0	58.5	95.4	59.6	59.3	67.6	58.5	66.3	90.5 \pm 4.6
walker2d-medium-v2	75.3	72.4	79.7	77.8	81.8	83.7	74.0	72.5	78.3	87.0 \pm 0.9
halfcheetah-medium-replay-v2	36.6	40.5	42.2	40.2	38.1	44.6	36.6	45.5	44.2	47.8 \pm 0.3
hopper-medium-replay-v2	18.1	37.2	96.8	93.6	97.5	60.9	82.7	95.0	94.7	101.3 \pm 0.6
walker2d-medium-replay-v2	26.0	27.0	61.2	49.8	49.5	81.8	66.6	77.2	73.9	95.5 \pm 1.5
halfcheetah-medium-expert-v2	55.2	42.8	79.8	53.3	93.4	90.7	86.8	91.6	86.7	96.8 \pm 0.3
hopper-medium-expert-v2	52.5	55.8	107.2	108.7	103.3	98.0	107.6	105.4	91.5	111.1 \pm 1.3
walker2d-medium-expert-v2	107.5	74.5	108.4	95.6	113.0	110.1	108.1	108.8	109.6	110.1 \pm 0.3
Average	51.9	50.1	75.3	72.9	76.1	75.3	74.7	77.6	77.0	88.0
AntMaze Tasks	BC	AWAC	BCQ	BEAR	Onestep RL	TD3+BC	DT	CQL	IQL	Diffusion-QL
antmaze-umaze-v0	54.6	56.7	78.9	73.0	64.3	78.6	59.2	74.0	87.5	93.4 \pm 3.4
antmaze-umaze-diverse-v0	45.6	49.3	55.0	61.0	60.7	71.4	53.0	84.0	62.2	66.2 \pm 8.6
antmaze-medium-play-v0	0.0	0.0	0.0	0.0	0.3	10.6	0.0	61.2	71.2	76.6 \pm 10.8
antmaze-medium-diverse-v0	0.0	0.7	0.0	8.0	0.0	3.0	0.0	53.7	70.0	78.6 \pm 10.3
antmaze-large-play-v0	0.0	0.0	6.7	0.0	0.0	0.2	0.0	15.8	39.6	46.4 \pm 8.3
antmaze-large-diverse-v0	0.0	1.0	2.2	0.0	0.0	0.0	0.0	14.9	47.5	56.6 \pm 7.6
Average	16.7	18.0	23.8	23.7	20.9	27.3	18.7	50.6	63.0	69.6
Adroit Tasks	BC	SAC	BCQ	BEAR	BRAC-p	BRAC-v	REM	CQL	IQL	Diffusion-QL
pen-human-v1	25.8	4.3	68.9	-1.0	8.1	0.6	5.4	35.2	71.5	72.8 \pm 9.6
pen-cloned-v1	38.3	-0.8	44.0	26.5	1.6	-2.5	-1.0	27.2	37.3	57.3 \pm 11.9
Average	32.1	1.8	56.5	12.8	4.9	-1.0	2.2	31.2	54.4	65.1
Kitchen Tasks	BC	SAC	BCQ	BEAR	BRAC-p	BRAC-v	AWR	CQL	IQL	Diffusion-QL
kitchen-complete-v0	33.8	15.0	8.1	0.0	0.0	0.0	0.0	43.8	62.5	84.0 \pm 7.4
kitchen-partial-v0	33.8	0.0	18.9	13.1	0.0	0.0	15.4	49.8	46.3	60.5 \pm 6.9
kitchen-mixed-v0	47.5	2.5	8.1	47.2	0.0	0.0	10.6	51.0	51.0	62.6 \pm 5.1
Average	38.4	5.8	11.7	20.1	0.0	0.0	8.7	48.2	53.3	69.0

Q&A